



THE UNIVERSITY of EDINBURGH
informatics

Operating Systems (INFR10079) 2023/2024 Semester 2

Introduction (Operating Systems and Computing Systems History)

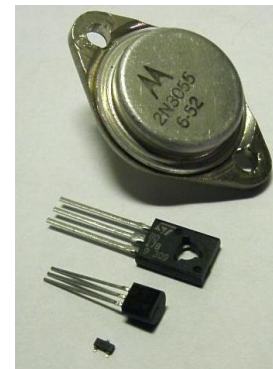
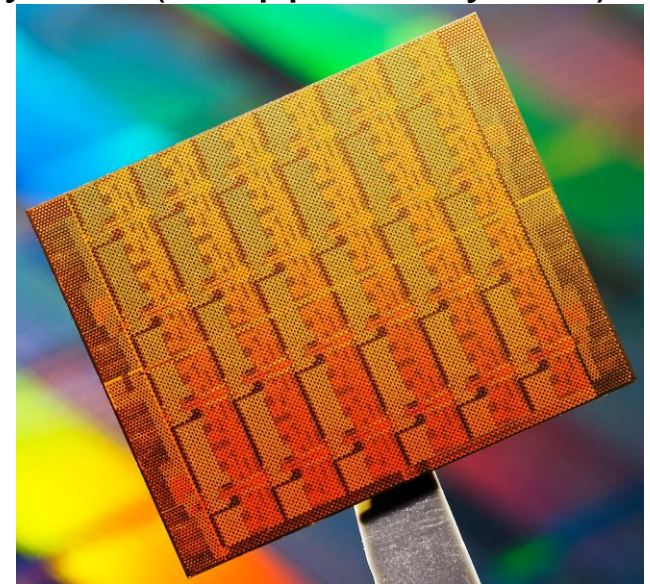
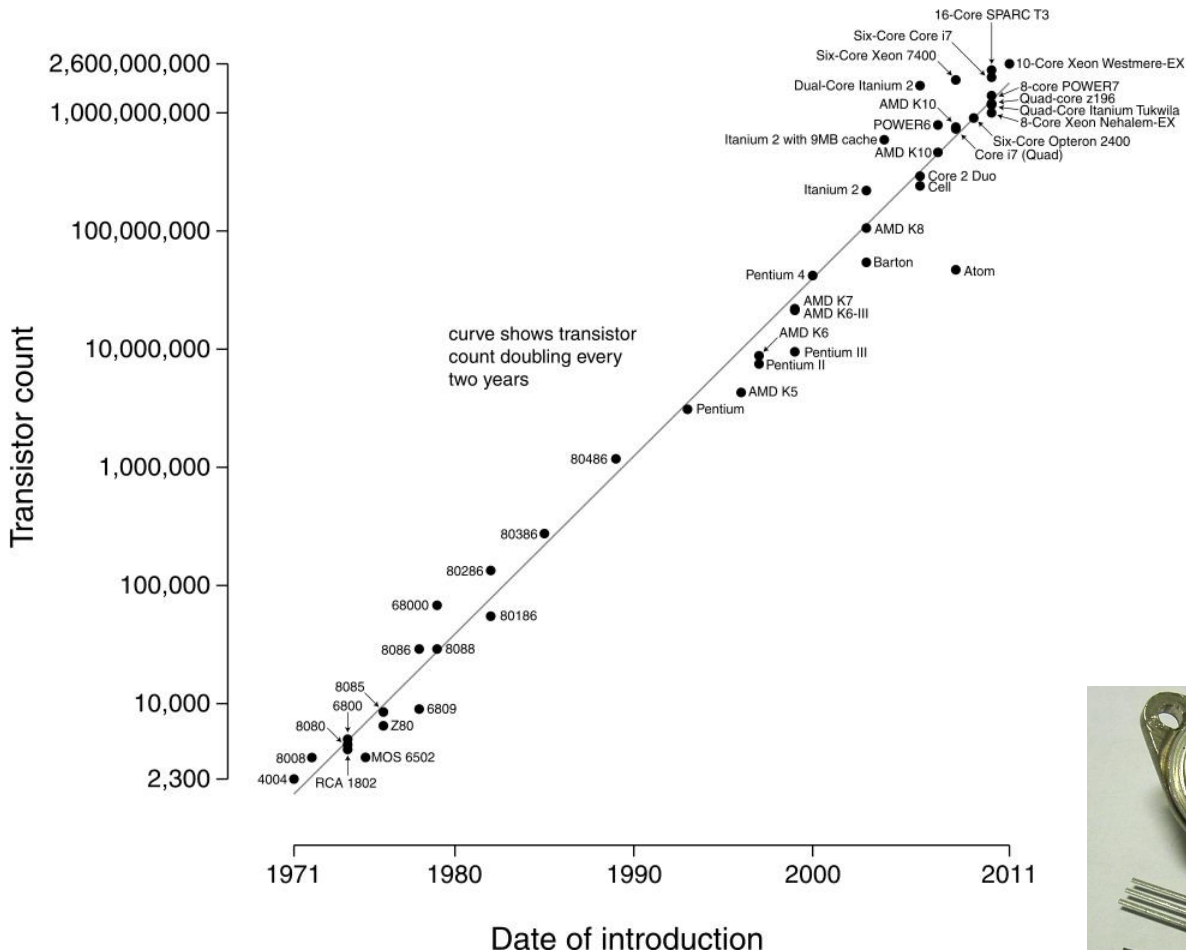
abarbala@inf.ed.ac.uk

Chapter 1.4.1, 1.5, 1.6, 1.7, 1.8, 1.10, Appendix A

Hardware Complexity Increases

Microprocessor Transistor Counts 1971-2011 & Moore's Law

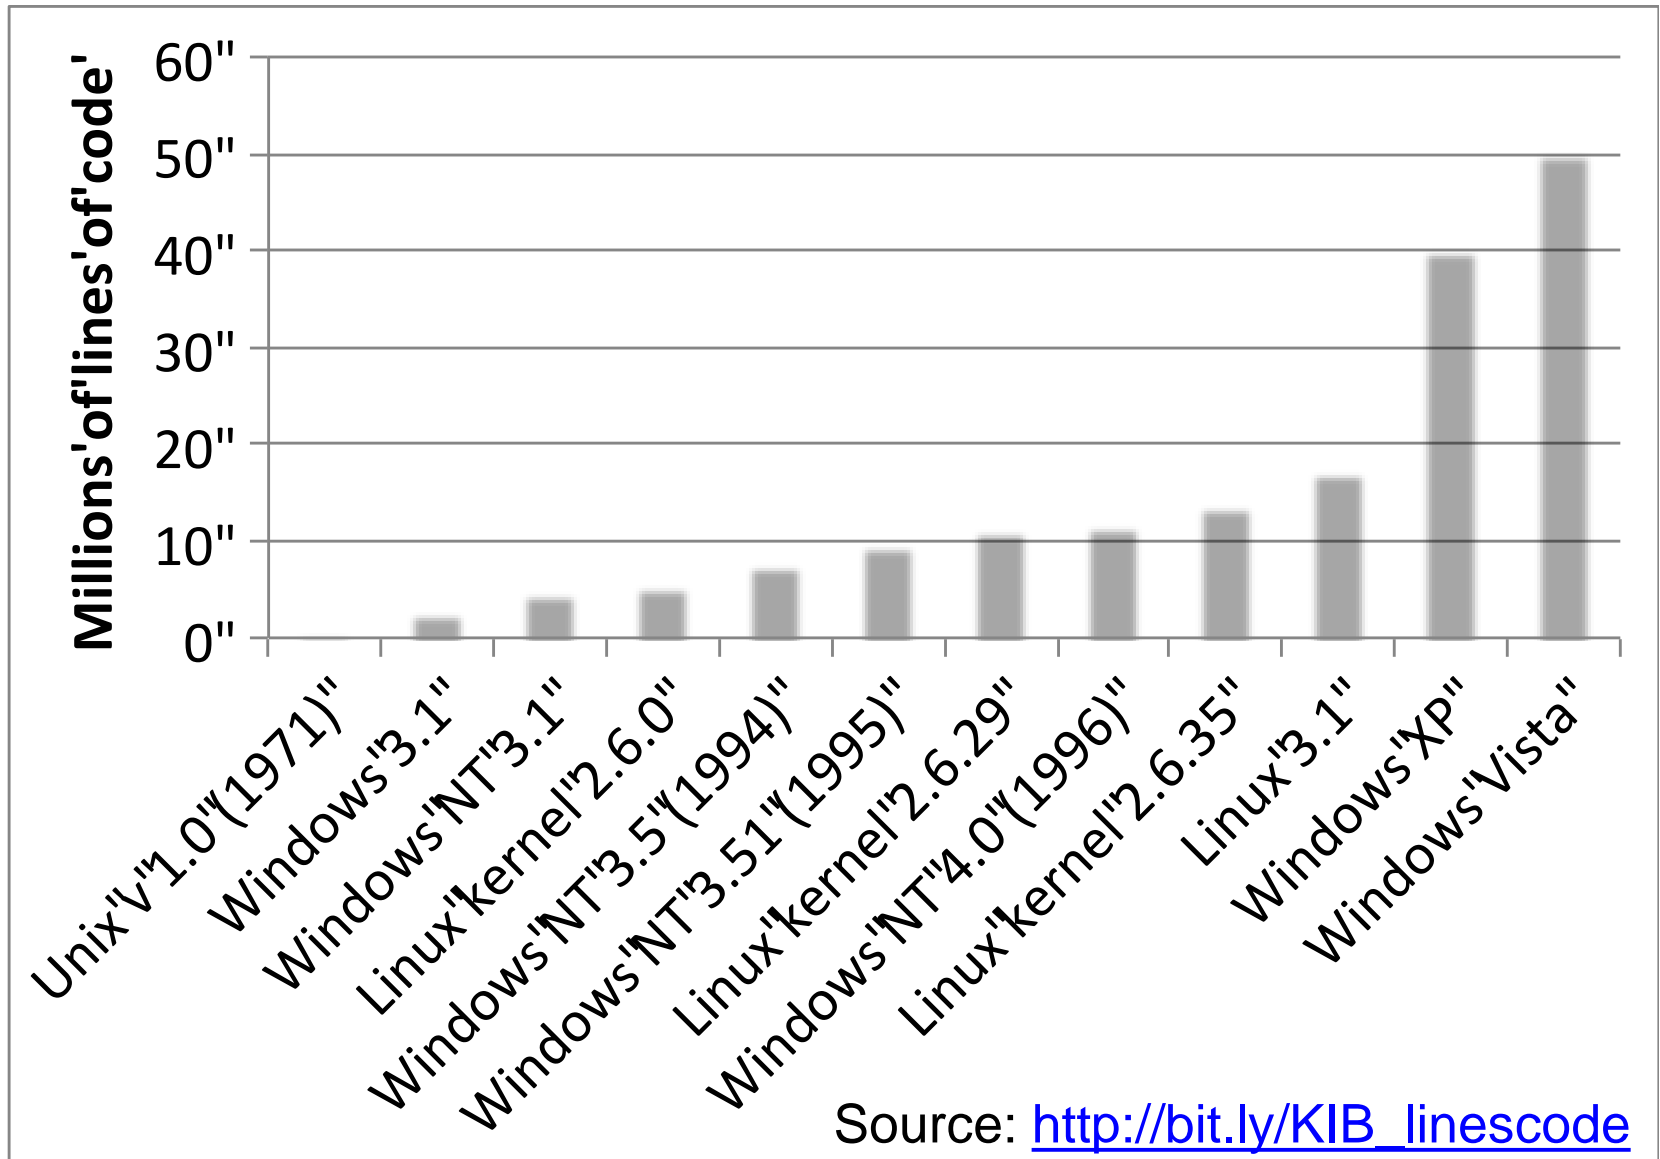
Moore's Law: 2X
transistors/Chip Every 1.5
years (or approx. 2 years)



Transistor: the
basic switching (0/1)
component (discrete)

https://en.wikipedia.org/wiki/Transistor_count
https://en.wikipedia.org/wiki/Moore%27s_law

Software Complexity Increases



Wolverhampton Instrument for Teaching Computing from Harwell (WITCH)



1951
No OS!

Run a
specific
calculation



Hardware Changes with Time

- **1960s:** mainframe computers (IBM)
- **1970s:** minicomputers (DEC)
- **1980s:** microprocessors and workstations (SUN), local-area networking, the Internet
- **1990s:** PCs (rise of Microsoft, Intel, Dell), the Web
- **2000s:**
 - Internet Services / Clusters (Amazon)
 - General Cloud Computing (Google, Amazon, Microsoft)
 - Mobile/ubiquitous/embedded computing (iPod, iPhone, iPad, Android)
- **2010s:** sensor networks, “data-intensive computing,” computers and the physical world
- **2020s:** it’s up to you!!

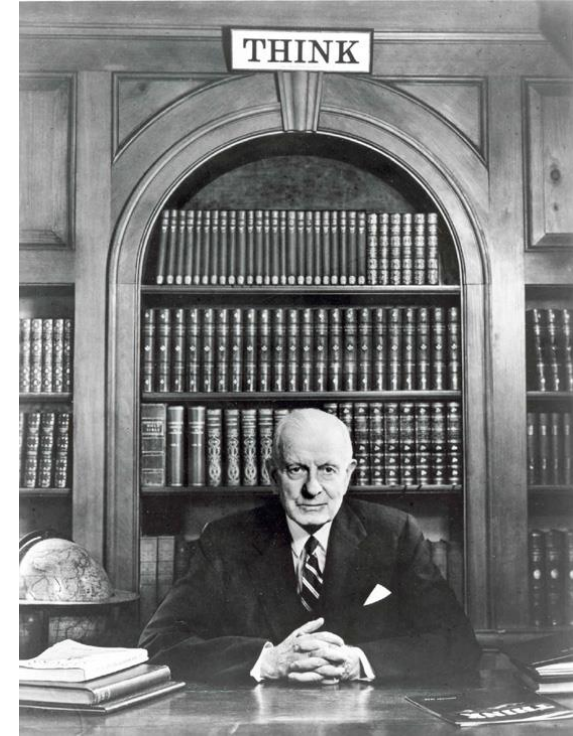


An OS History Lesson

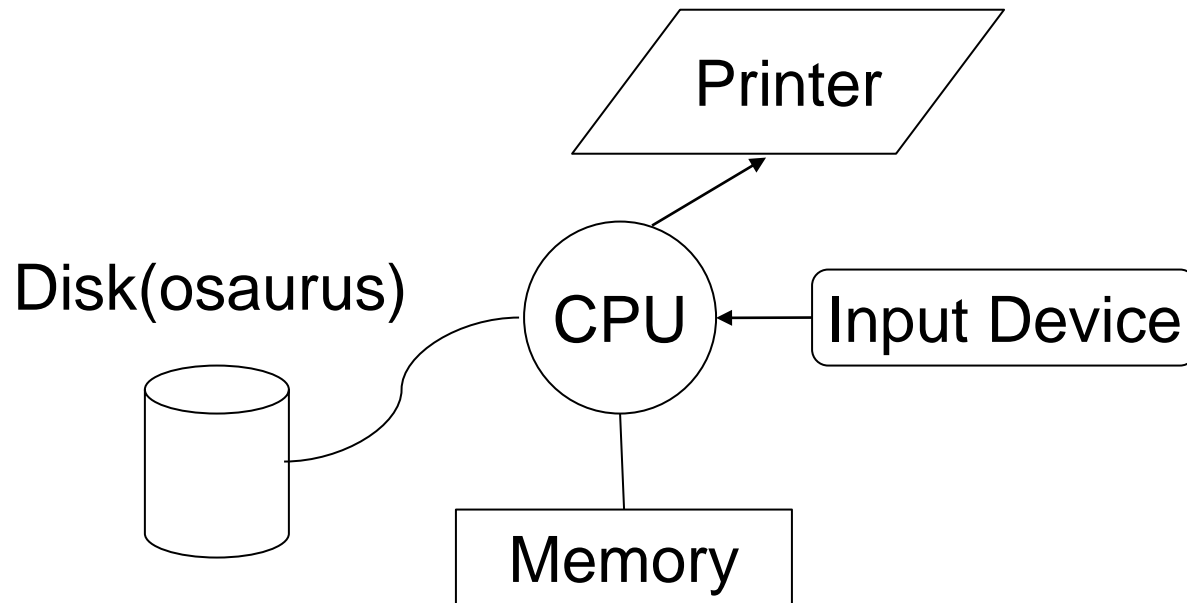
- Operating systems
 - Almost as old as computers
 - Result of a **60 year long** evolutionary process
- Why looking at their history? (see next slides)
 - Learn about their evolution
 - Born out of several needs
 - Clarify what some of their functions are, and why
- Are they still evolving?
 - **Yes!** (not the topic of this course)

At the Beginning...

- 1943
 - T.J. Watson (created IBM):
“I think there is a world market for maybe five computers.”
- Fast forward ... 1950
 - There are maybe 20 computers in the world
 - Execute one program at the time (no OS)
 - Unbelievably expensive
 - Imagine: machine time is more valuable than person time!
 - Ergo: **efficient use of the hardware** is paramount
 - Operating systems are born
 - They carry with them the vestiges of these ancient forces



Primordial Computer



The OS as a Linked Library

One
program at
the time

One user at
the time

- At the very beginning...
 - OS as **a library of code** that you **linked into** your program
 - Programs were **loaded in their entirety into memory, and executed**
 - “OS” had an “API” that let you control
 - Disk
 - Printer
 - ...
 - Human interfaces were literally **switches** and blinking lights
 - When done running program, leave and turn the computer over to next person

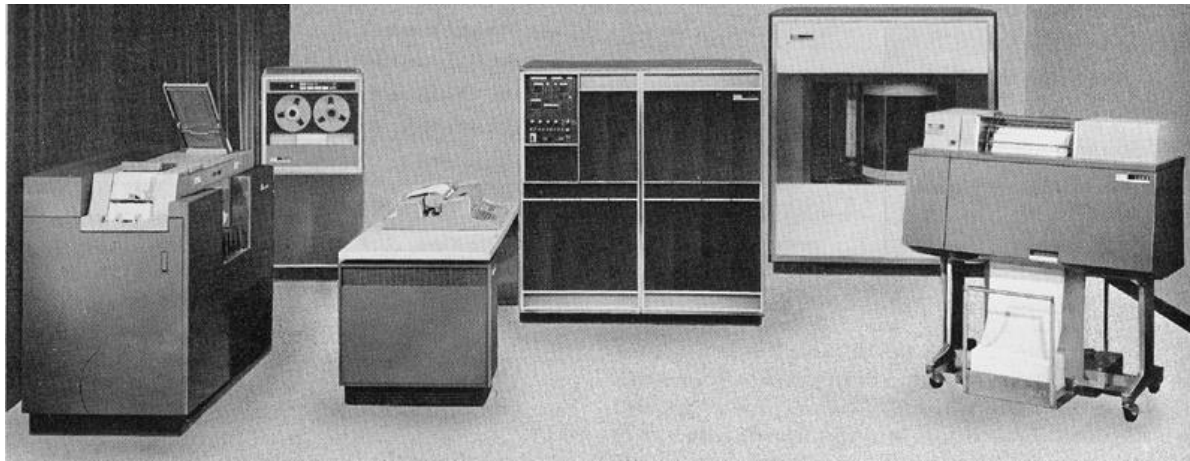


Asynchronous I/O

One
program at
the time

One user at
the time

- The disk(osaurus) was really slow
- Add hardware so disk could operate without tying up the CPU
 - **Disk controller**
- Hot shot programmers could now write code that
 - Starts an I/O
 - Goes off and does some computing
 - Checks if the I/O is done at some later time
- **Upside**
 - Helps increase (expensive) CPU utilization
- **Downsides**
 - It's hard to get right
 - The benefits are job specific



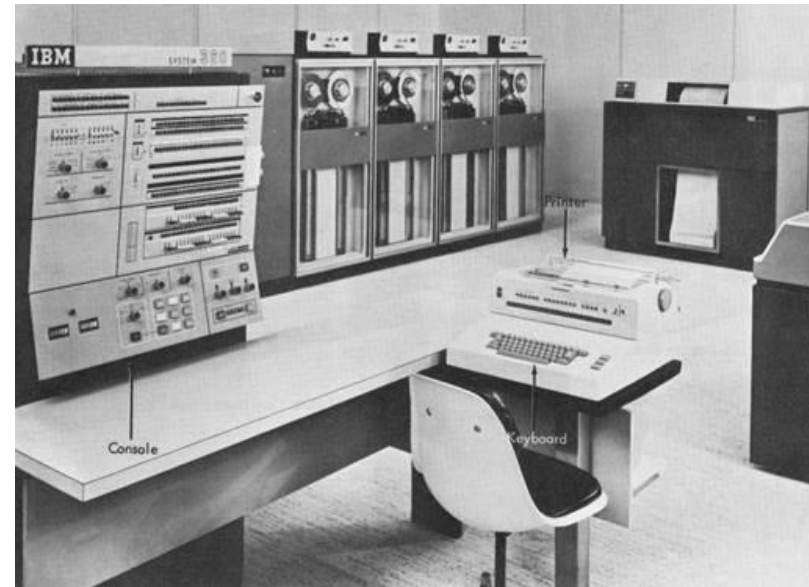
IBM 1401

Multiprogramming

Multiple
programs at
the time

One user at
the time

- To further increase system utilization:
multiprogramming OSs
 - Multiple runnable jobs in memory at once
 - Overlaps I/O of one job with computing of another
 - While one job waits for I/O completion, another job uses the CPU
 - Can get rid of asynchronous I/O within individual jobs
 - Life of application programmer becomes simpler
 - Only the OS programmer needs to deal with asynchronous events
 - How to tell when devices are done?
 - Interrupts
 - Polling



IBM System 360

Timesharing

Multiple
programs at
the time

Multiple
users at the
time

- To support interactive use, create a **timesharing OS**
 - multiple terminals for one machine
 - each user has illusion of entire machine to him/herself
 - optimize response time, perhaps at the cost of throughput
- Timeslicing
 - divide CPU equally among the users
 - if job is truly interactive (e.g., editor), then can jump between programs and users faster than users can generate load
 - permits users to interactively view, edit, debug running programs
- MIT CTSS system (operational 1961) among the first timesharing systems
 - only one user memory-resident at a time (32KB memory!)
- MIT Multics system (operational 1968) first large timeshared system
 - nearly all OS concepts can be traced back to Multics!

Multiple Terminals

- In early 1980s, a single timeshared VAX-11/780 ran computing for all of CSE
- A typical VAX-11/780 was 1 MIPS (1 MHz) and had 1MB of RAM and 100MB of disk
- An Apple iPhone 4 is 1GHz (x1000), has 512MB of RAM (x512), and 32GB of flash (x320)

Multiple
programs at
the time

Multiple
users at the
time



Parallel Systems

Targeting
High
Performance

- Some applications can be written as multiple parallel **threads** or **processes**
 - can speed up the execution by running multiple threads/processes simultaneously on multiple CPUs [Burroughs D825, 1962]
 - need OS and language primitives for dividing program into multiple parallel activities
 - need OS primitives for fast communication among activities
 - degree of speedup dictated by communication/computation ratio
 - many flavors of parallel computers today
 - SMPs (symmetric multi-processors)
 - MPPs (massively parallel processors)
 - NOWs (networks of workstations)
 - Massive clusters (Google, Amazon.com, Microsoft)
 - Computational grid (SETI @home)

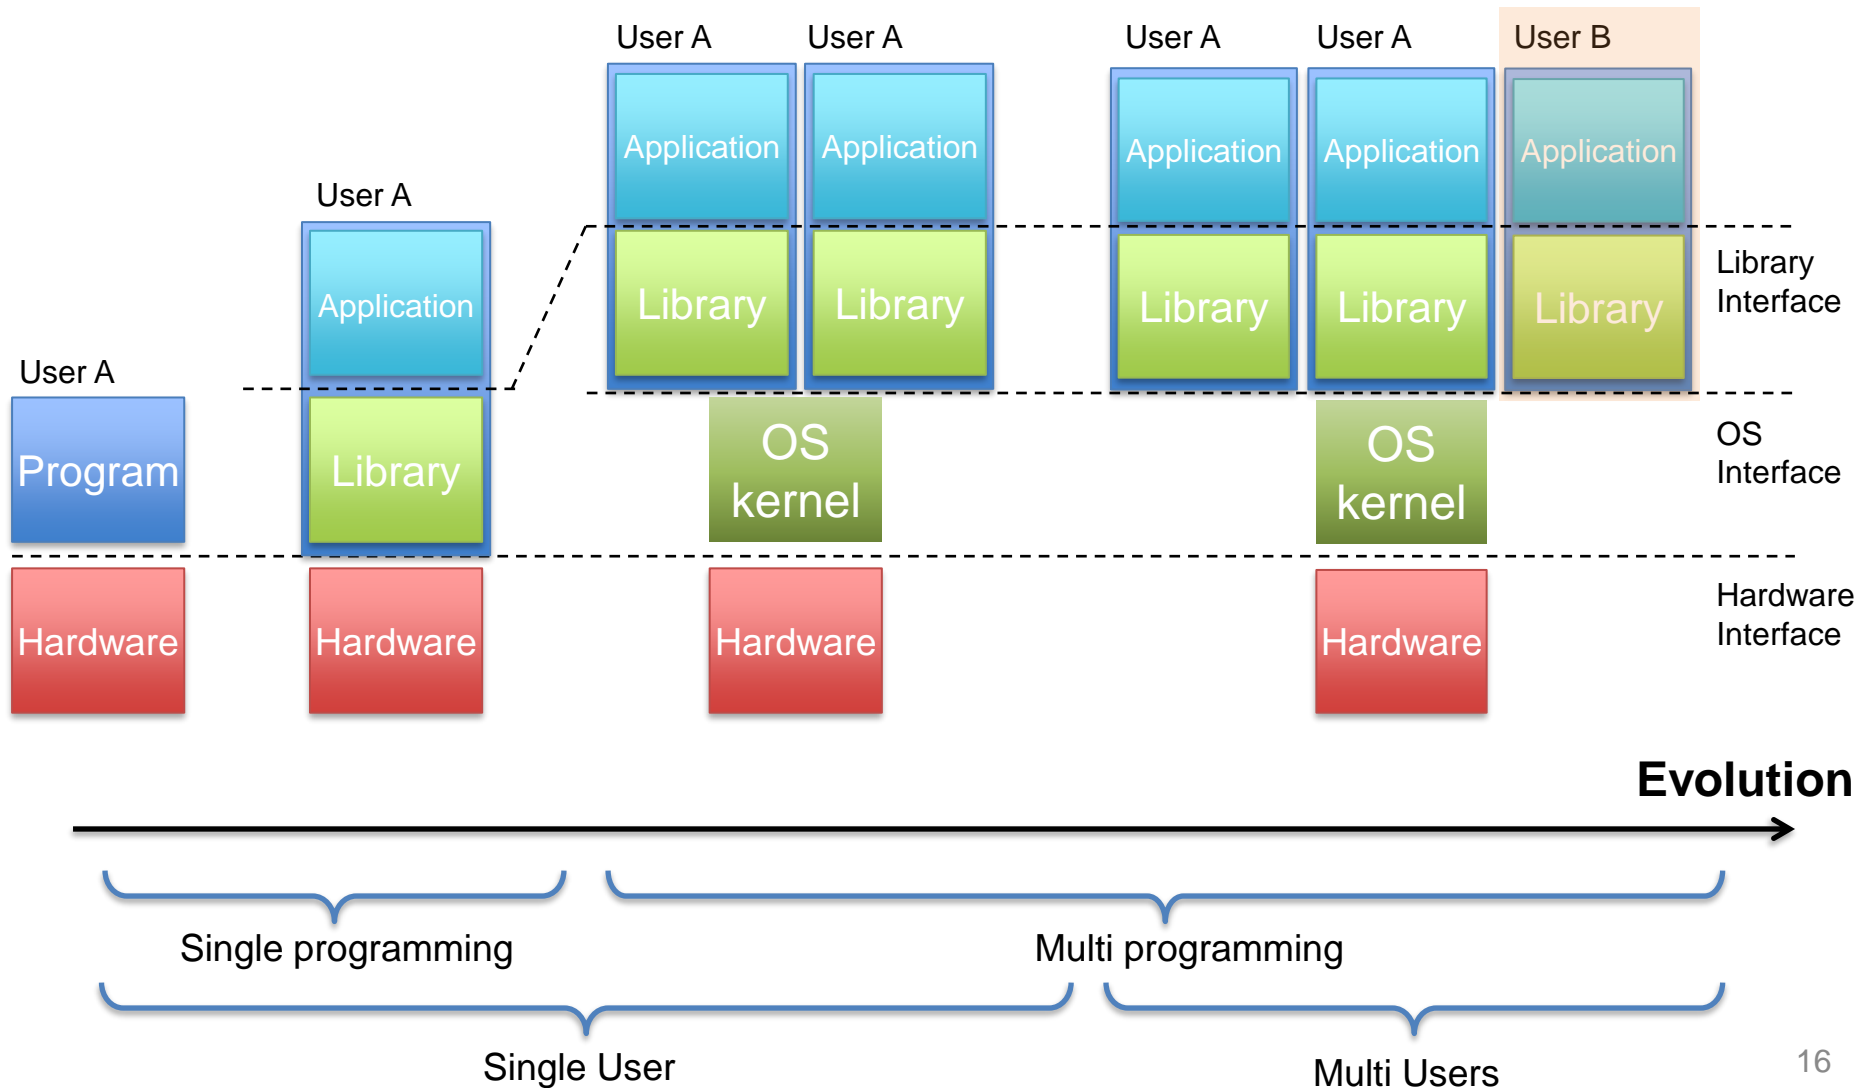
Personal Computing

Targeting
Personal Use

- Primary goal was to enable new kinds of applications
- Bit mapped display [Xerox Alto, 1973]
 - new classes of applications
 - new input device (the mouse)
- Move computing near the display
 - why?
- Window systems
 - the display as a managed resource
- Local area networks (Ethernet)
 - why?

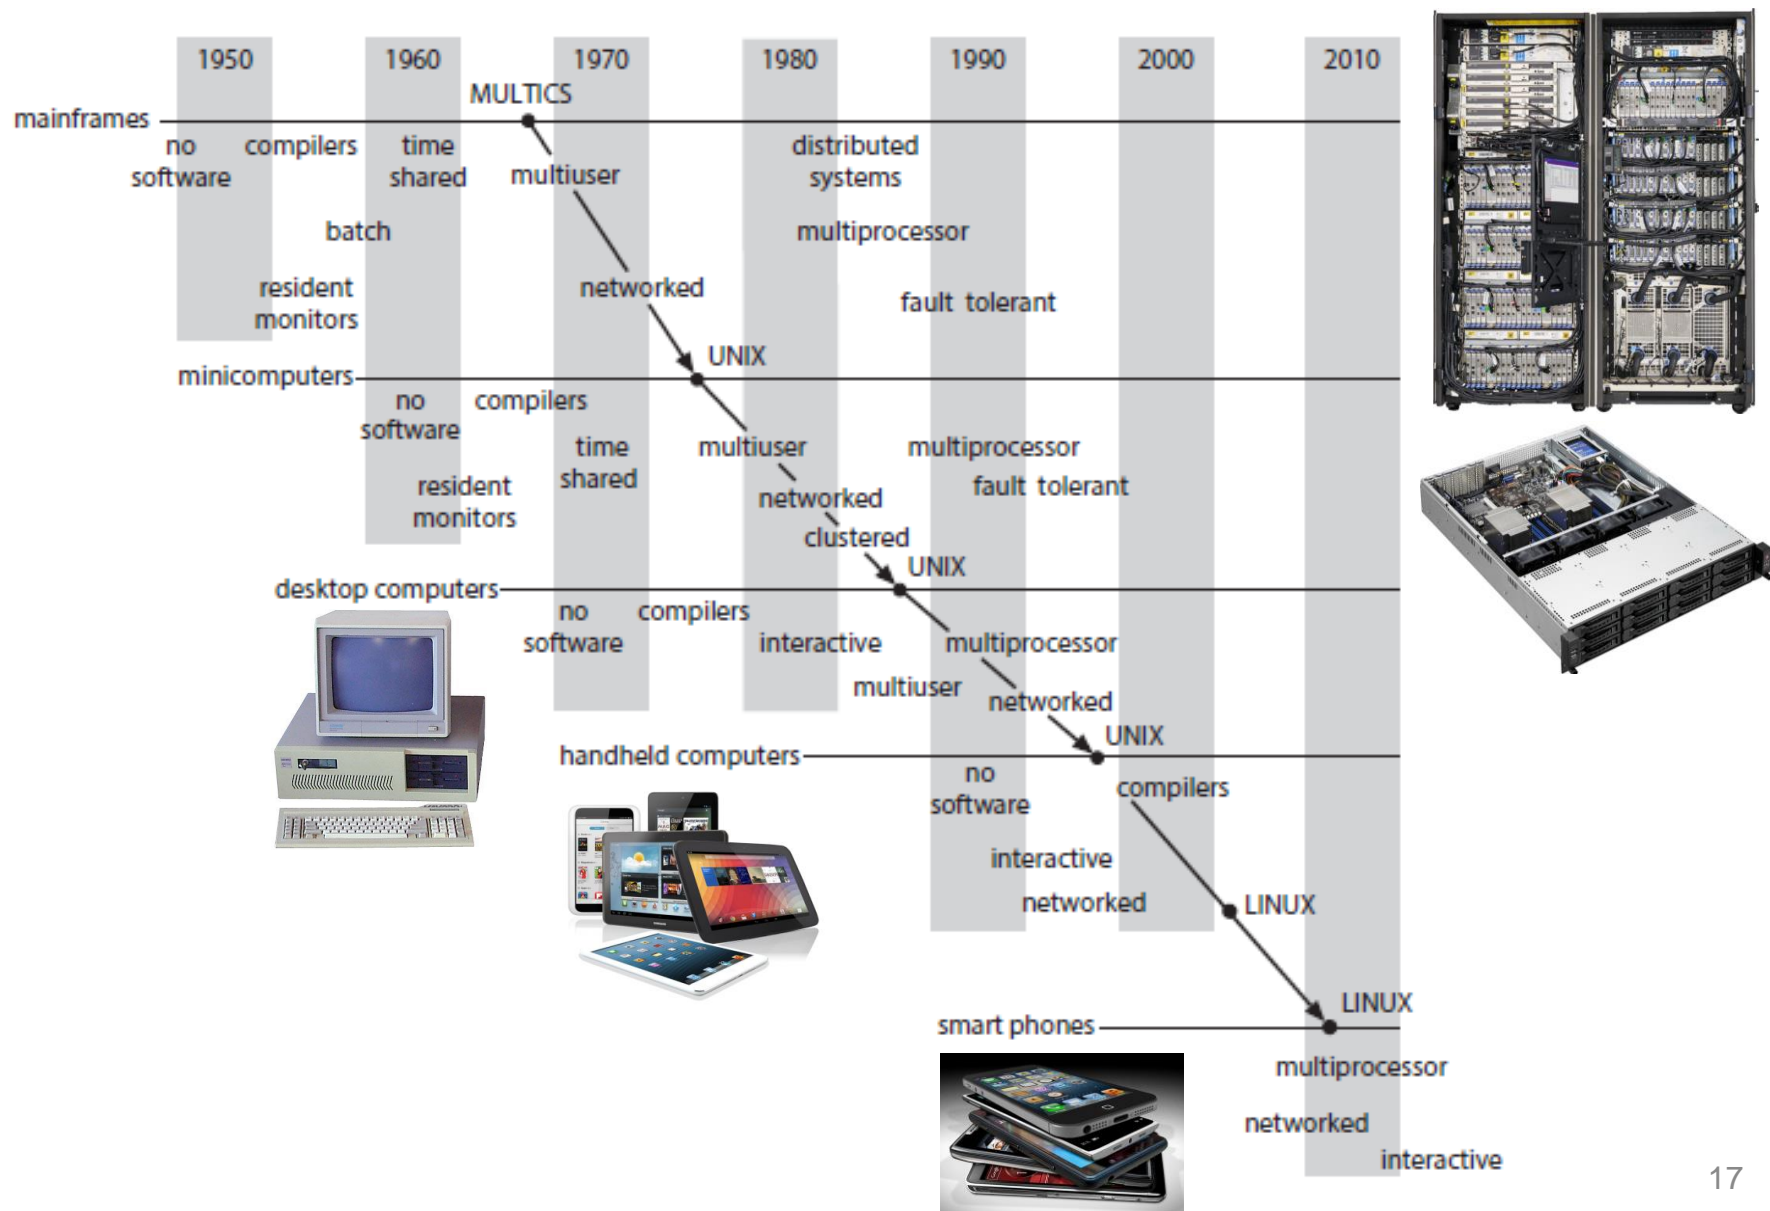


Progression of Programming Abstractions



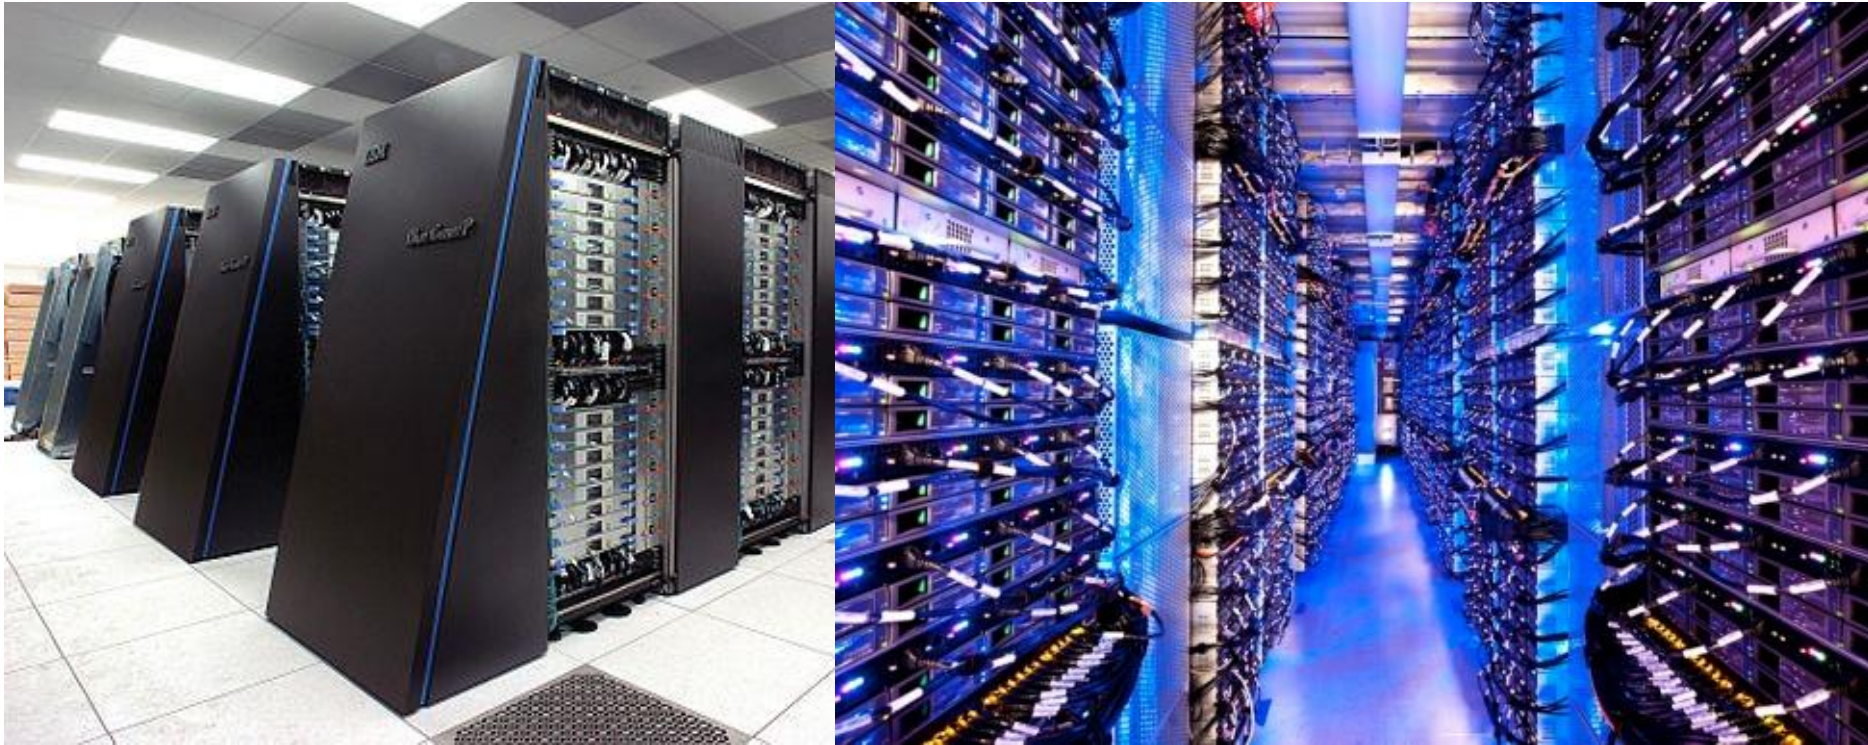
Single Computer

Progression of Concepts and Form Factors



When a Single Computer is Not Enough

- Interconnect multiple computers together
- Supercomputers, clusters, data-center ...

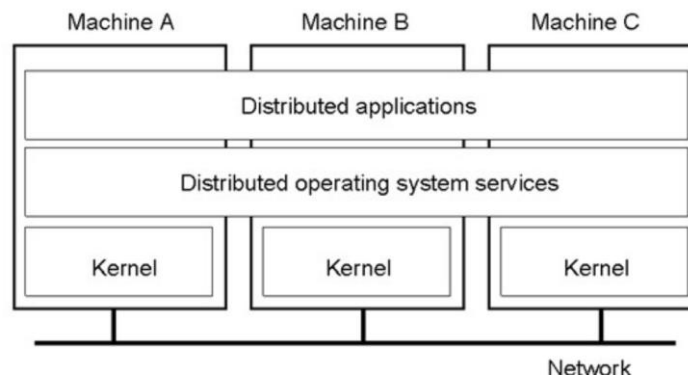


Interconnected Computers

- Interconnection
 - Bus Extension
 - Networking
- Approaches
 - Distributed OS
 - Client/server Computing
 - Peer-to-Peer systems
 - Cloud Computing
 - ...

Distributed OS

- Distributed systems to facilitate use of distributed resources
 - workstations on a LAN
 - servers across the Internet
- Supports applications running among multiple computers
 - interprocess communication
 - message passing, shared memory
 - networking stacks
- Sharing of distributed resources (hardware, software)
 - load balancing, authentication and access control, ...



Client/Server Computing

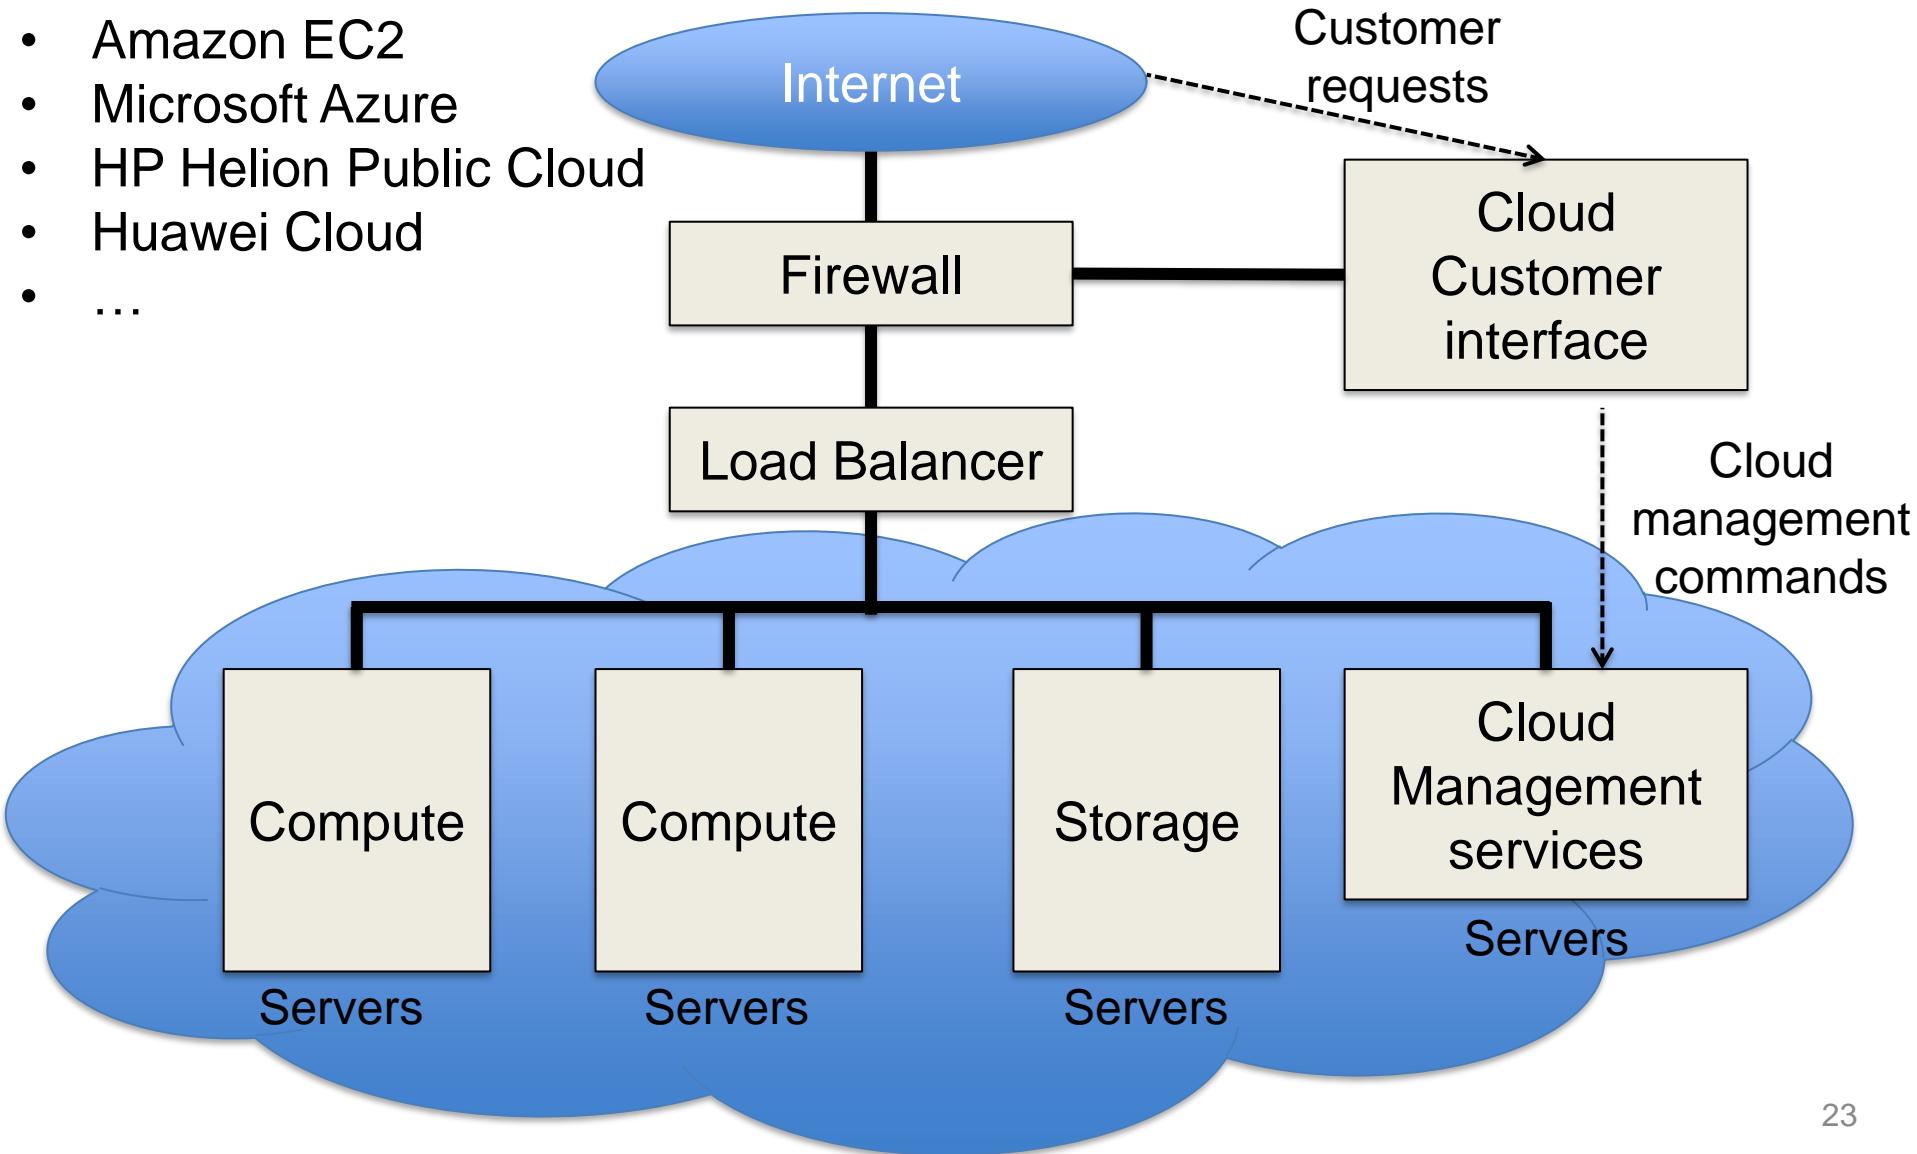
- Mail server/service
- File server/service
- Print server/service
- Compute server/service
- Game server/service
- Music server/service
- Web server/service
- etc.

Peer-to-Peer (p2p) Systems

- Napster
- Gnutella
- BitTorrent
 - example technical challenge: self-organizing overlay network
 - technical advantage of BitTorrent?
 - er ... legal advantage of BitTorrent?

Cloud Computing

- Amazon EC2
- Microsoft Azure
- HP Helion Public Cloud
- Huawei Cloud
- ...



The major OS issues

- **Structure:** how is the OS organized?
- **Sharing:** how are resources shared across users?
- **Naming:** how are resources named (by users or programs)?
- **Security:** how is the integrity of the OS and its resources ensured?
- **Protection:** how is one user/program protected from another?
- **Performance:** how do we make it all go fast?
- **Reliability:** what happens if something goes wrong (either with hardware or with a program)?
- **Extensibility:** can we add new features?
- **Communication:** how do programs exchange information, including across a network?

More OS issues...

- **Concurrency:** how are parallel activities (computation and I/O) created and controlled?
- **Scale:** what happens as demands or resources increase?
- **Persistence:** how do you make data last longer than program executions?
- **Distribution:** how do multiple computers interact with each other?
- **Accounting:** how do we keep track of resource usage, and perhaps charge for it?

There are tradeoffs, not right and wrong!

Why Should One Learn Operating Systems?

- You may not ever build an OS
- But almost all code runs on top an OS
 - Almost every computing device runs an OS
- Hence, you need to understand the foundations
 - As a Computer Scientist or Computer Engineer
- Knowledge of how OSes work is crucial to proper, efficient, effective, and secure **programming**

It is a Great Time to Study Operating Systems!

- Many open-source OS projects
 - Check, study, and modify the code
 - Help find and fix bugs
- Emulators to run and debug OSes
 - Existent OSes
 - Historical/Future OSes (hardware not available)
- Renewed interest in OSes from major IT Companies
 - Rise of unikernels
 - Fully formally verified OSes (e.g., Huawei Harmony)
 - New OSes (e.g., Google Fuchsia)