



THE UNIVERSITY of EDINBURGH
informatics

Operating Systems (INFR10079) 2022/2023 Semester 2

Processes (Interprocess Communication)


abarbala@inf.ed.ac.uk

Chapter 3.4, 3.5, 3.6, 3.7, 3.8, excluding: 3.7.2, 3.7.3, 3.8.2.1
4.6.2, 20.9.1, C3.3


Interprocess Communication

- **Independent** processes
- **Cooperating** processes
 - Information **sharing**
 - More applications interested in the same information
 - Computation **speedup**
 - To exploit parallel hardware
 - **Modularity**
 - Reusability of components
- Require **interprocess communication** (IPC) mechanism to send and receive data
 - Shared memory
 - Message passing

Interprocess Communication




Shared memory



Message passing

Shared Memory



- Allow processes to communicate and synchronize
 - **Sharing** part of address space
 - **OS doesn't mediate** communication (no overhead)
 - Usually, **OS prevents** processes to access each other memory
 - Processes should agree (have permission) to void this restriction
- Data
 - Format **decided by application**
 - Direct access (not mediated by the OS) – **very fast**
 - Application programmer fully manages the data transfer – **not trivial**
- Possible use cases
 - Passing of large (single) objects (C&P, image, table, etc.)
 - Notification variable

Example of Shared Memory Code (POSIX)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main() {
    const int SIZE = 4096; /* the size (bytes) of shared memory object */
    const char *name = "OS"; /* name of the shared memory object */
    const char *message 0 = "Hello"; /* written to shared memory */
    const char *message 1 = "World!"; /* written to shared memory */
    int fd; /* shared memory file descriptor */
    char *ptr; /* pointer to shared memory object */

    /* create the shared memory object */
    fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    /* configure the size of the shared memory object */
    ftruncate(fd, SIZE);
    /* memory map the shared memory object */
    ptr = (char *)
        mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* write to the shared memory object */
    sprintf(ptr, "%s", message 0);
    ptr += strlen(message 0);
    sprintf(ptr, "%s", message 1);
    ptr += strlen(message 1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main() {
    const int SIZE = 4096; /* the size (bytes) of shared memory object */
    const char *name = "OS"; /* name of the shared memory object */

    int fd; /* shared memory file descriptor */
    char *ptr; /* pointer to shared memory object */

    /* open the shared memory object */
    fd = shm_open(name, O_RDONLY, 0666);
    /* memory map the shared memory object */
    ptr = (char *)
        mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* read from the shared memory object */
    printf("%s", (char *)ptr);
    /* remove the shared memory object */
    shm_unlink(name);


    return 0;
}
```



Try at
home!

Message Passing

- Allow processes to communicate and synchronize
 - **Without sharing** part of address space
 - **OS mediates** communication (introduce overhead)
- Processes on the **same** machine and **among different** inter-networked machines
 - Not possible with shared memory
- Message-passing facility provides at least two operations
 - **send(message)**
 - **receive(message)**
- Communication link
 - Several **implementation tradeoffs**, e.g., messages size
 - Fixed
 - Variable



Message Passing – Naming

- Communicating processes must **refer to each other**
- **Direct** communication
 - **Symmetric:** Explicit name of sender and receiver
 - `send(P, message)` – send message to process P
 - `receive(Q, message)` – receive message from process Q
 - **Asymmetric:** Explicit at least one end
 - `send(p, message)` – send to p
 - `receive(id, message)` – receive from any process, sender saved in id
- **Indirect** communication
 - No need to know/explicit in advance sender and/or receiver
 - Mailboxes (e.g., POSIX Mailbox)
 - `send(A, message)` – into mailbox A
 - `receive(A, message)` – from mailbox A
 - A mailbox can be accessed by more than two processes
 - Between processes multiple mailboxes may exist



Naming is a
problem
also in
shared
memory

Message Passing – Synchronization

- Different **design options** to implement `send()` / `receive()`
 - **Blocking** or **synchronous**
 - **Nonblocking** or **asynchronous**
- Different **combinations** may be offered
 - Blocking send
 - Nonblocking send
 - Blocking receive
 - Nonblocking receive
- **Rendezvous**
 - When both `send()` and `receive()` **are blocking**
 - Simple solution to the consumer-producer problem***

Message Passing – Buffering

- Message exchanged reside in **temporary buffers/queues**
 - **Zero** capacity (no buffering)
 - No message waiting in it
 - Sender must block until the recipient receives the message
 - **Bounded** capacity
 - Finite length n ; at most n messages can reside in it
 - If the queue is not full when a new message is sent, that is placed in the queue and the sender can continue execution
 - If the link is full, the sender has to wait
 - **Unbounded** capacity
 - Infinite queue
 - The sender never blocks

Example of message passing: Pipes

- A pipe acts as a **conduit** allowing two processes to communicate
 - One way data flow
- **Anonymous** (ordinary)
 - Between parent and child
- **Named**
 - between any pair of processes
 - In UNIX, FIFO




Figure 3.20 File descriptors for an ordinary pipe.

Example of Pipe Code (POSIX)

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#define READ_END 0
#define WRITE_END 1
int main(void) {
    char write_msg[256] = "Greetings";
    char read_msg[256];
    int fd[2];
    pid_t pid;

    if (pipe(fd) == -1) { /* create the pipe */
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    pid = fork(); /* fork a child process */
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    if (pid > 0) { /* parent process */
        close(fd[READ_END]); /* close the unused end of the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1); /* write to the pipe */
        close(fd[WRITE_END]); /* close the write end of the pipe */
    }
    else { /* child process */
        close(fd[WRITE_END]); /* close the unused end of the pipe */
        read(fd[READ_END], read_msg, 256); /* read from the pipe */
        printf("read %s", read_msg);
        close(fd[READ_END]); /* close the read end of the pipe */
    }
    return 0;
}
```



Try at
home!

Client-Server Communication

- **Sockets** Abstraction
 - Endpoint for communication
 - Identified by an **IP address** concatenated with a **port number**
 - Servers implementing specific services (SSH, FTP, and HTTP) listen to well-known ports
 - an SSH server listens to port 22; an FTP server listens to port 21; and a web, or HTTP, server listens to port 80
- **Remote Procedure Call (RPC)**
 - Abstract the **procedure-call mechanism**
 - For use between systems with network connections
 - **Similar** in many respects **to the IPC**
 - Uses message-based communication to provide remote service

Signals

- **OS mechanism to notify a process** (one way)
 - Doesn't carry information
- From **the OS**
 - Can be thought as a **software-generated interrupt/exception**
 - Synchronous – e.g., division by zero error
 - Asynchronous – e.g., timer/alarm
 - **Syscall: process -> OS; Signal: OS -> process**
- From **other processes** (including the process itself)
 - Only notification, no data **communication method**
 - Management – e.g., kill a process
 - Synchronization – e.g., POSIX RT Signals
 - etc.

Handling Signals (Linux Example)

- **Signal handler**
 - Code to process a signal
 - No return value (void)
- Handler must be registered
 - Or OS default action
- Send a signal
 - `kill()`
 - `raise()`

https://www.cs.auckland.ac.nz/references/unix/digital/APS33DTE/DOCU_006.HTM#realtime-handler-sec

```
#include <unistd.h>
#include <signal.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
#define SIG_STOP_CHILD SIGUSR1
```

```
main() {
    pid_t pid;
    sigset_t newmask, oldmask;

    if ((pid = fork()) == 0) { /* Child */
        struct sigaction action;
        void catchit();
        sigemptyset(&newmask);
        sigaddset(&newmask, SIG_STOP_CHILD);
        sigprocmask(SIG_BLOCK, &newmask, &oldmask);
        action.sa_flags = 0;
        action.sa_handler = catchit;
        if (sigaction(SIG_STOP_CHILD, &action, NULL) == -1) {
            perror("sigusr: sigaction");
            _exit(1);
        }
        sigsuspend(&oldmask);
    }
    else { /* Parent */
        int stat;
        sleep(10);
        kill(pid, SIG_STOP_CHILD);
        pid = wait(&stat);
        printf("Child exit status = %d\n", WEXITSTATUS(stat));
        _exit(0);
    }
}

void catchit(int signo) { /* Signal Handler */
    printf("Signal %d received from parent\n", signo);
    _exit(0);
}
```

Try at
home!