



THE UNIVERSITY of EDINBURGH  
**informatics**

# **Operating Systems (INFR10079) 2023/2024 Semester 2**

## **Threads (Basics)**

[abarbala@inf.ed.ac.uk](mailto:abarbala@inf.ed.ac.uk)

Chapter 4.1 (all), 4.2 (all)

# Overview

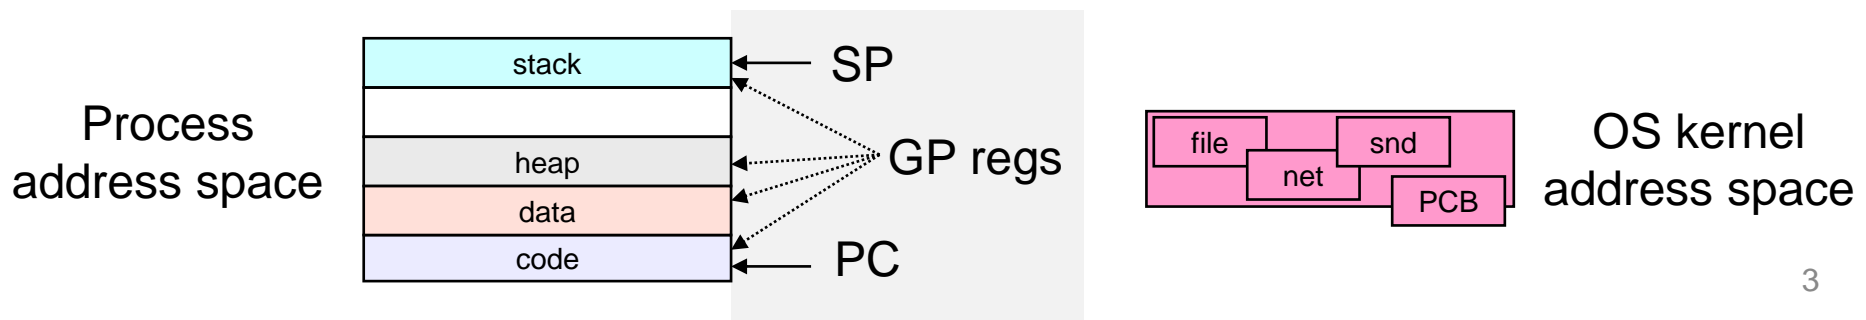
- Concurrency vs Parallelism
- Process vs Thread
- Threads

# What's “in” a process?

From  
previous  
slide-set

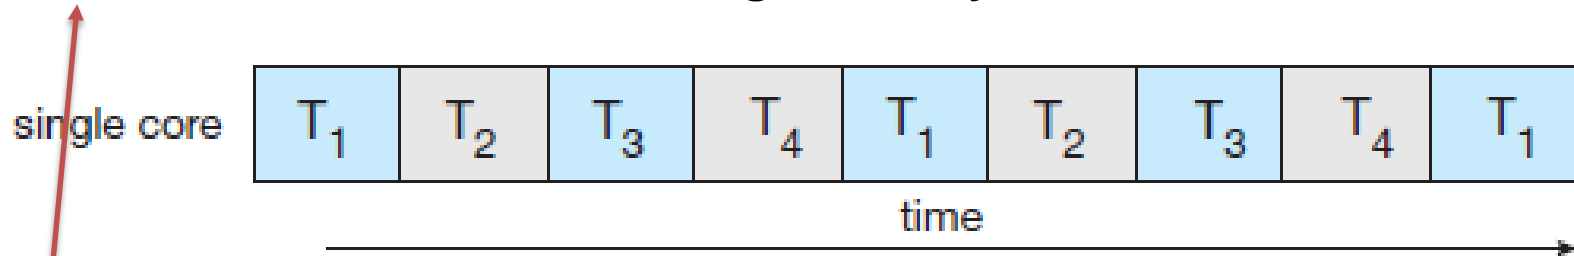
- A process consists of (at least)
  - An **address space**, containing
    - Code (instructions) for the running program
    - Data for the running program (static data, heap data, stack)
  - A **CPU state**, consisting of
    - Program counter (PC), indicating the next instruction
    - Stack pointer, current stack position
    - Other general-purpose register values
  - A set of **OS resources**
    - Open files, network connections, sound channels, ...

instruction  
flow

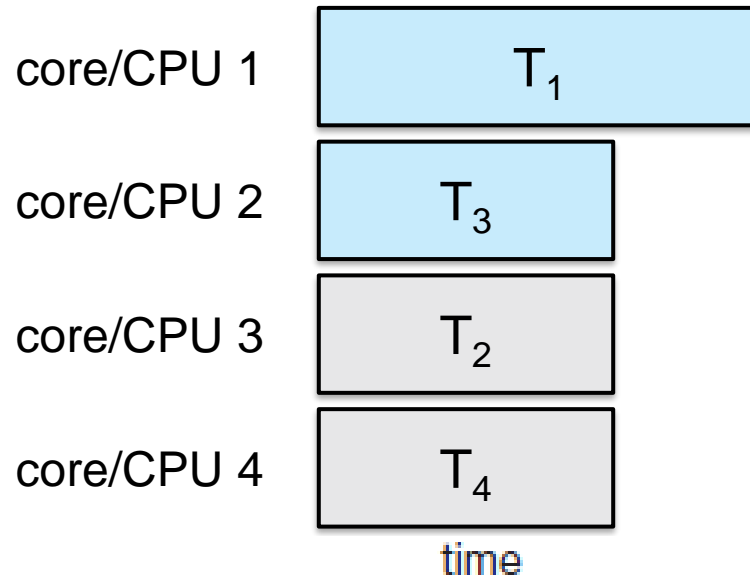


# Concurrency vs Parallelism


- ❑ **Multiple tasks:** Task 1 ( $T_1$ ), Task 2 ( $T_2$ ), Task 3 ( $T_3$ ), Task 4 ( $T_4$ )
- ❑ **Concurrent** execution on a **single-core system**



- ❑ **Parallel execution** on a **multicore/multiprocessor system**

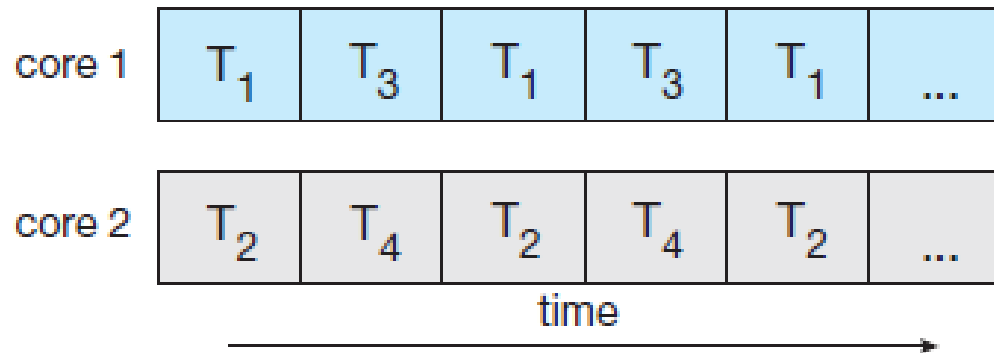


All tasks  
make  
progress at  
the “same  
time”



# Concurrency and Parallelism



- Both **Concurrency** and **Parallelism**



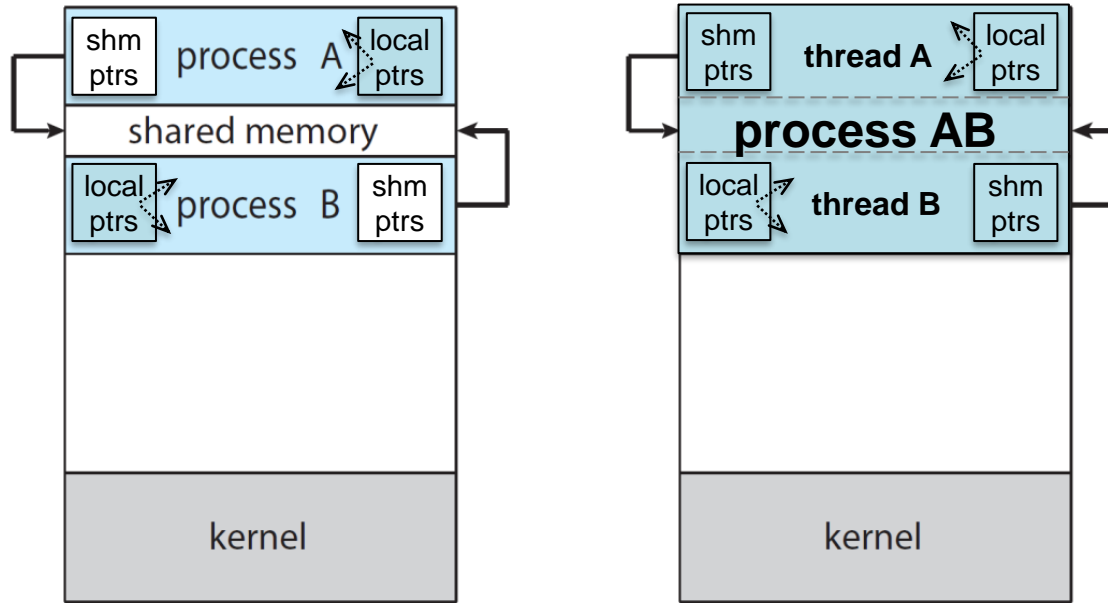
- Multiple **processes** to get *concurrency* and *parallelism*
  - Programs (code) of distinct processes are isolated from each other
- What if they need to **communicate/share data**?*
  - Message passing, OS in the middle – **slow**
  - What about **Shared memory**?

# Concurrent/parallel Communicating Processes

- Given the **process abstraction**
  1. Fork several processes
  2. Cause each of them to *map* to the **same** memory to share data
    - See `shmget()` API for one way to do this
  3. Make them both to *open* the **same** OS resources
- Cumbersome
  - OS resources not shared by default
- Limited shareability
  - Not all pointers work
- Inefficient
  - **Space**: PCB, page tables, etc.
  - **Time**: creating OS structures, fork/copy address space, etc.




# From Processes to Threads



- Multiple **processes** to get *concurrency* and *parallelism*
  - Programs (code) of distinct processes are isolated from each other
- Multiple **threads** to get *concurrency* and *parallelism*
  - Threads “**share a process**” – same address space, OS resources
  - Threads **different instruction flows** – private stack, CPU state

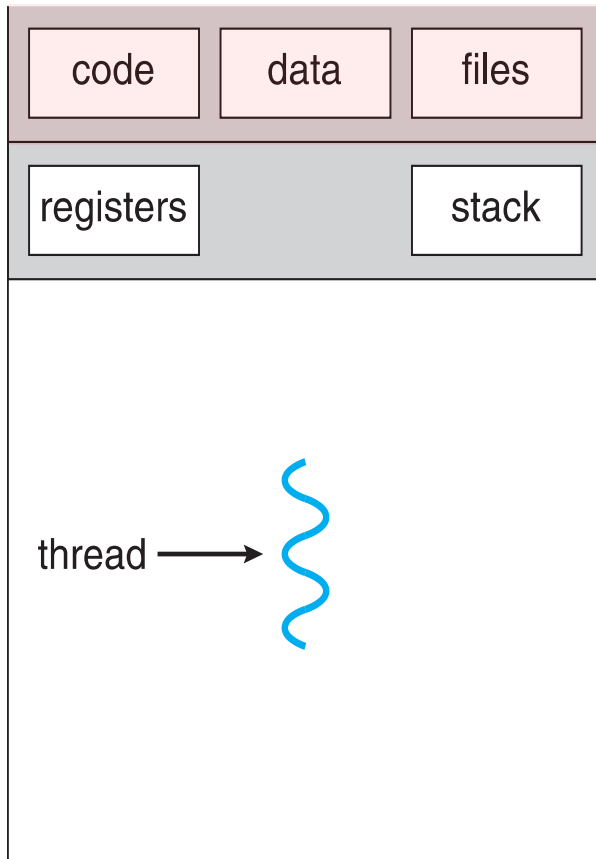
# Threads

- **Key idea**
  - **Separate** the foundational components of a *process* (address space, execution state, OS resources)
  - Into different **abstractions/entities**
    - **PROCESS:** address space, OS resources
    - **THREAD:** CPU state (execution state)
      - program counter, stack pointer, other registers

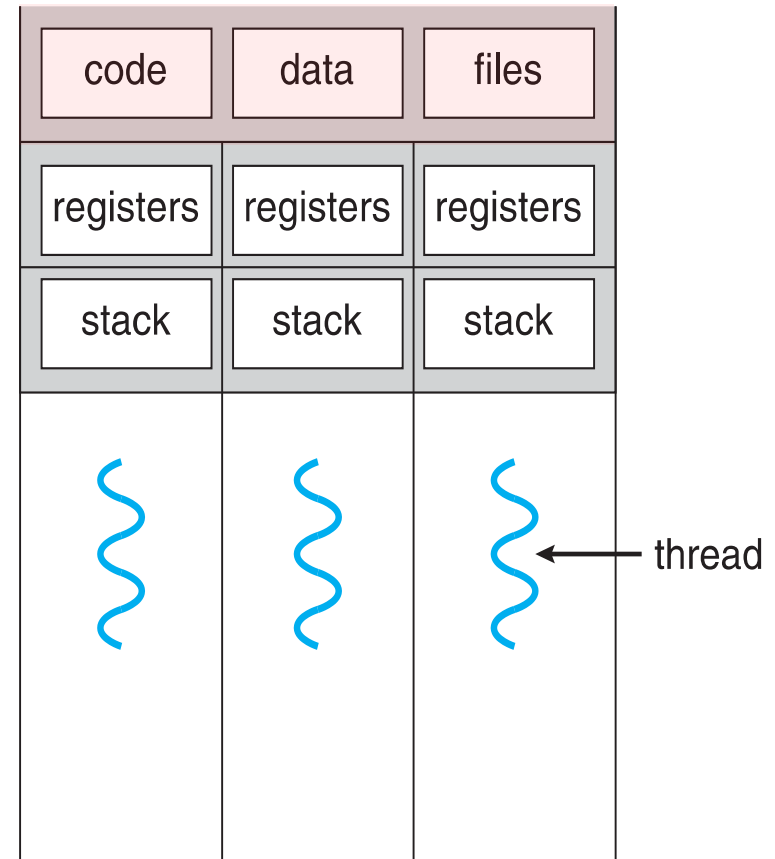




# Single-threaded and Multithreaded Processes



single-threaded process




multithreaded process

# Use Case Scenario

- Various **instruction flows**
  - run the *same or different code*
  - access the *same data* (or part of it)
  - have the *same privileges*
  - use the *same OS resources*
- Each **instruction flow** has **hardware execution state**
  - Execution stack and stack pointer (SP)
    - Traces state of procedure calls made
  - Program counter (PC)
    - Next instruction to be executed
  - Set of general-purpose processor registers and their values

# Threads and Processes



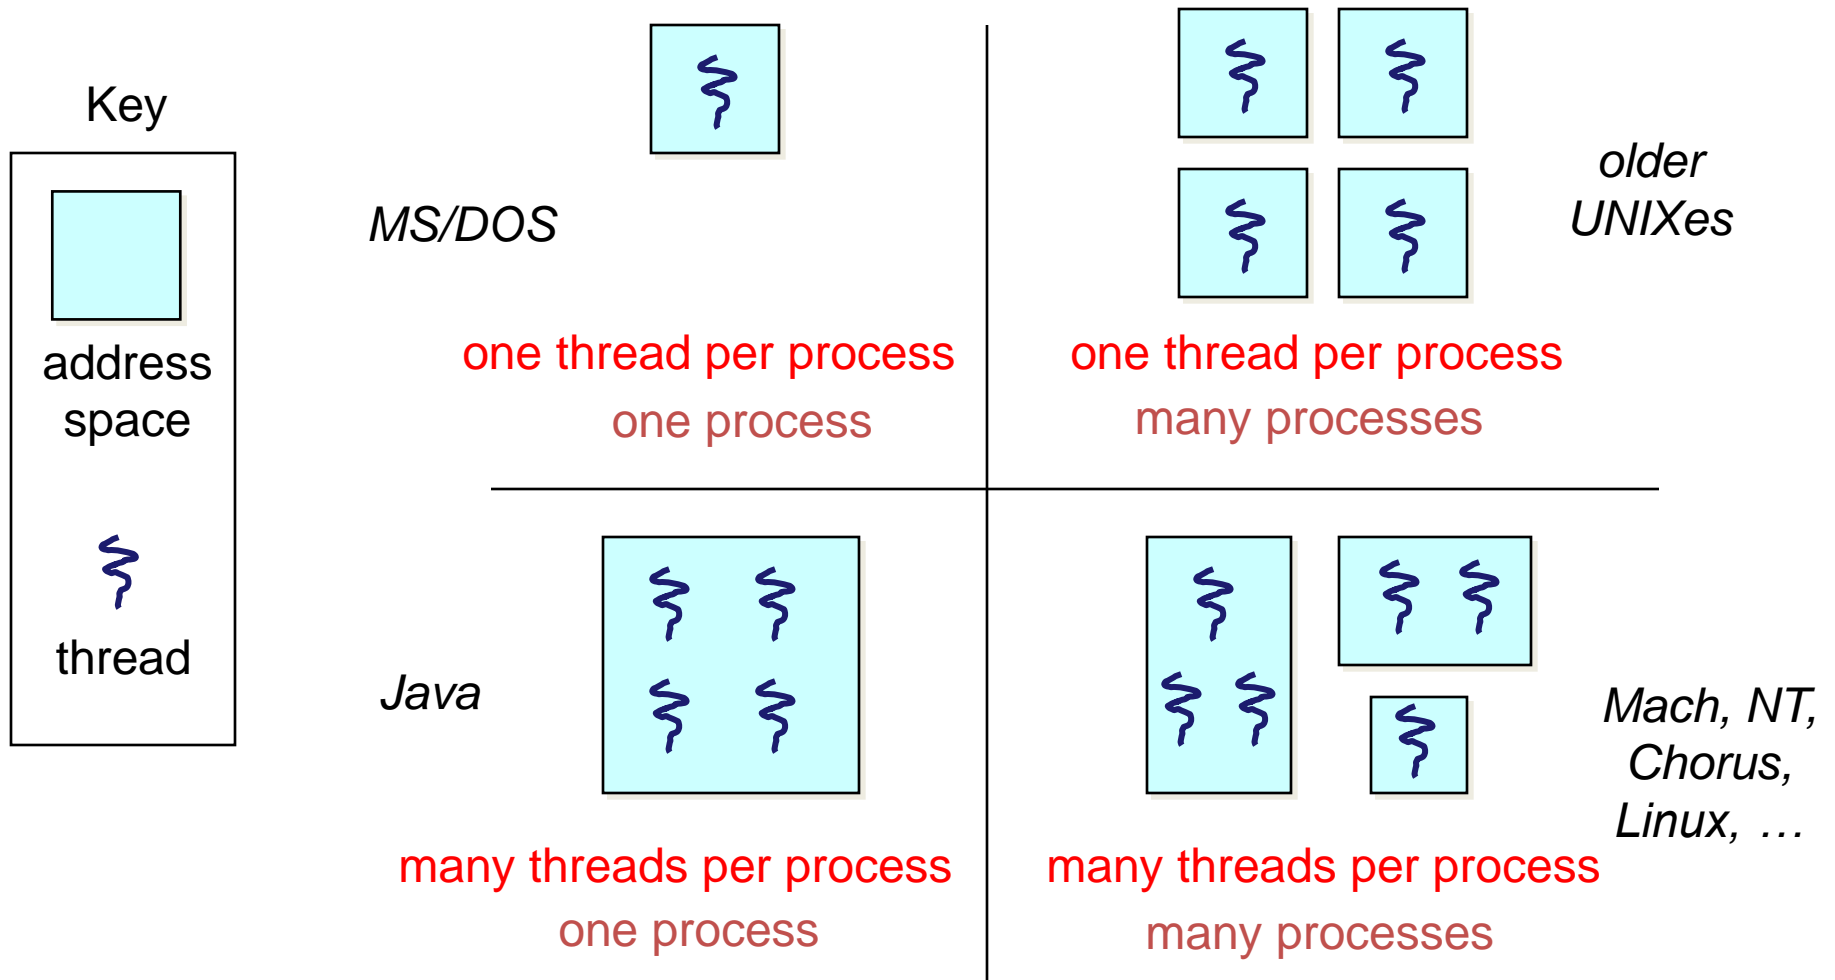
- Most modern OS's (Mach (Mac OS), Chorus, Windows, UNIX) support
  - **Process**: defines the address space and process' OS resources
  - **Thread**: defines a sequential **execution flow** within a process
- A **thread** is **bound to** a single process (thus address space)
  - However, processes (and address spaces) can have **multiple threads** executing within them
  - Sharing data between threads is **cheap**: all see the same address space
  - Creating threads is **cheap** too!
- **Threads** become the **unit of scheduling**
  - But depends on implementation (see next slides)
  - Processes are just **containers** in which threads execute

# Communication

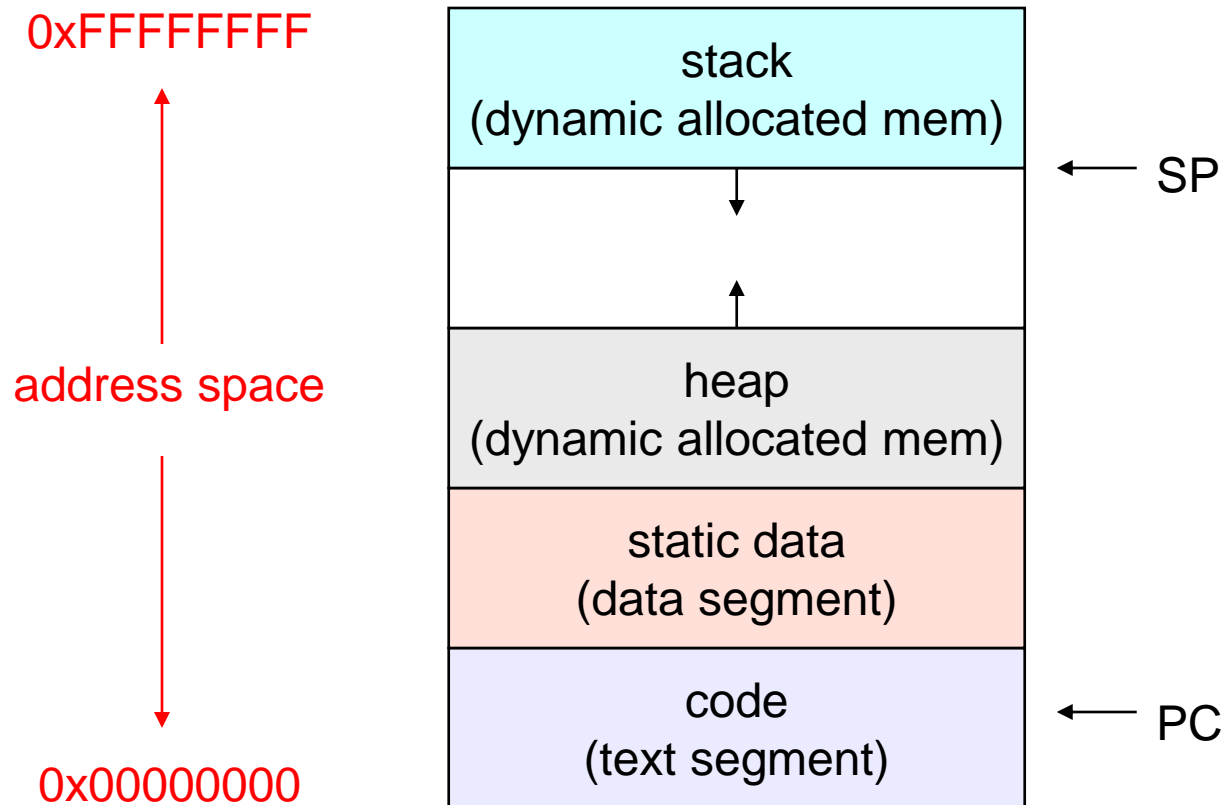
- **Threads** are diverse execution flows sharing an address space (and OS resources)
- Address spaces provide **isolation**
  - If you can't name it, you can't read nor write it
- **Threads** are in the same address space
  - Same name space (memory addresses)
  - Update a **shared variable**



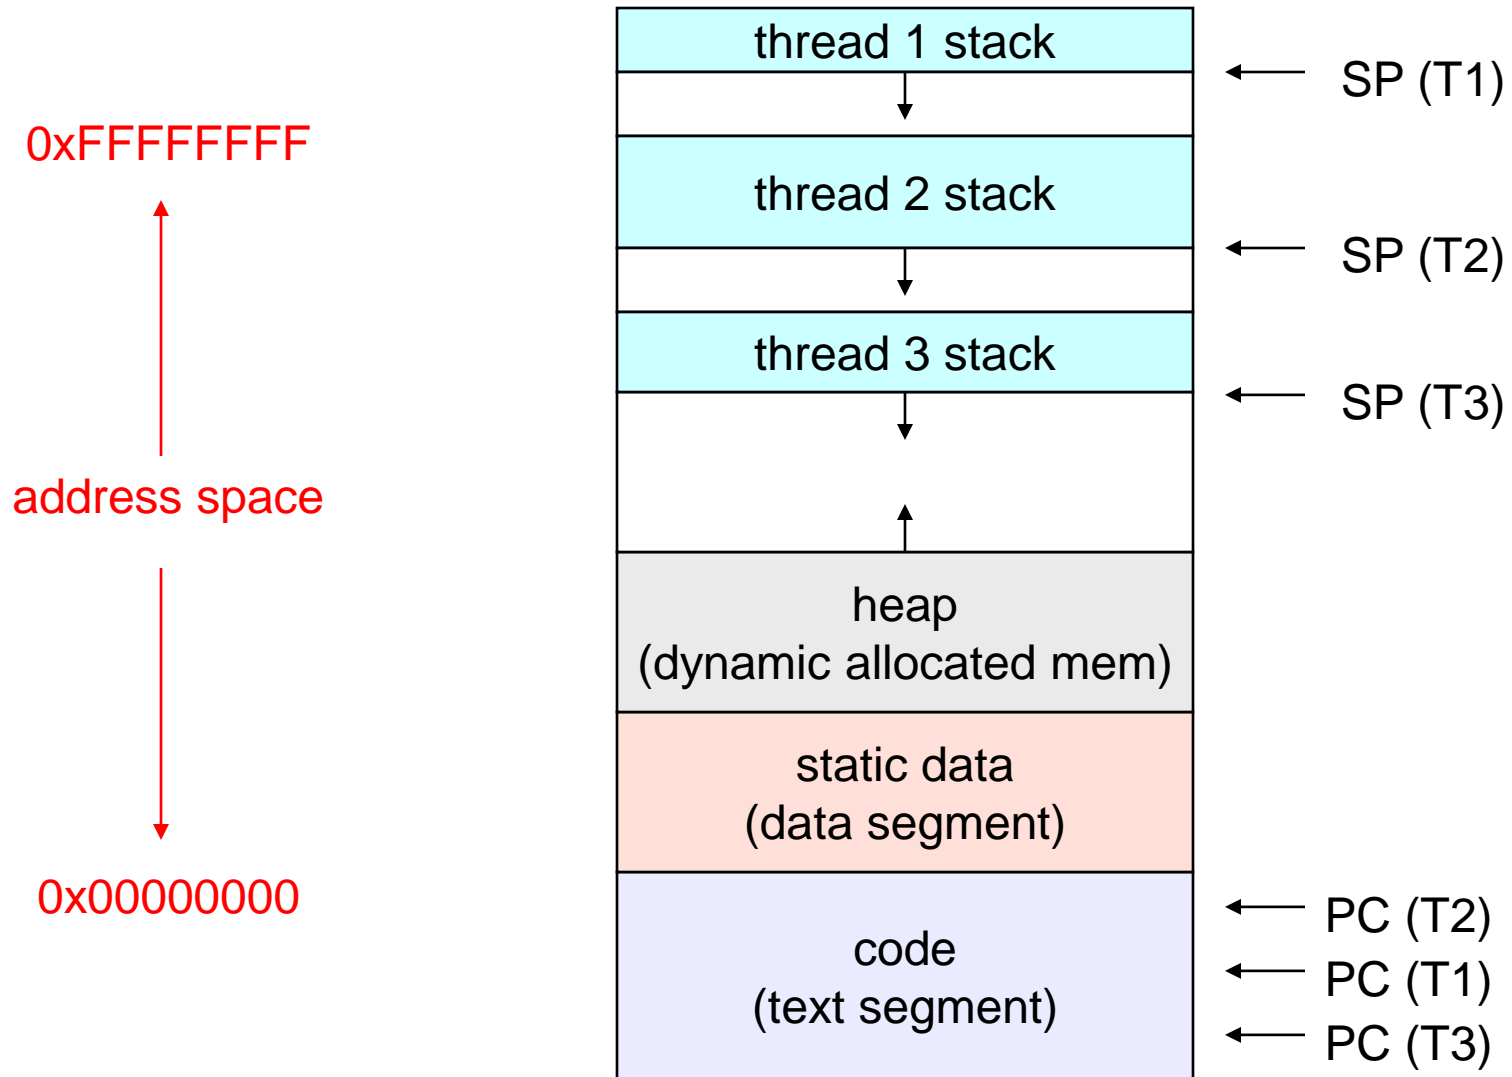
# Historical Design Space



# (Old) Process Address Space (32bit)

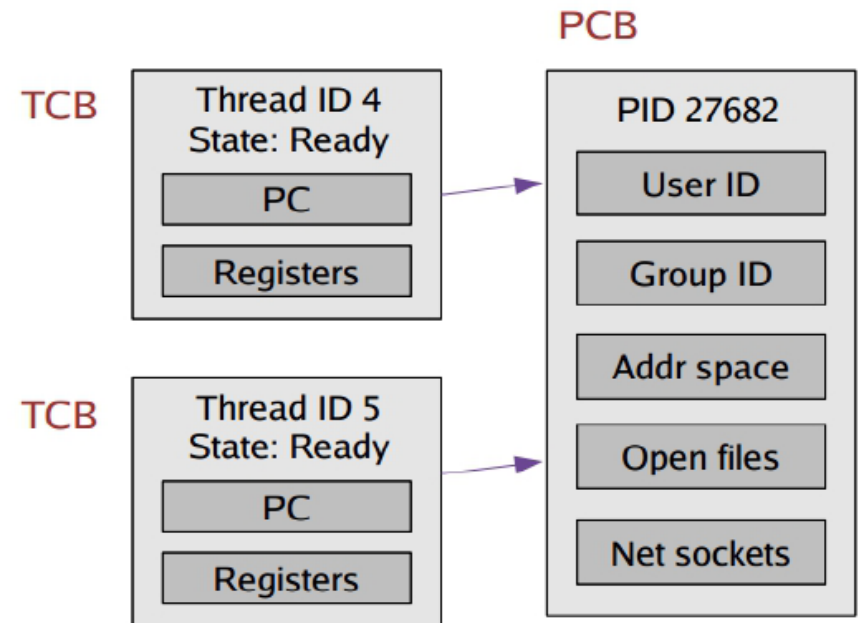


# (New) Address Space with **Threads** (32bit)



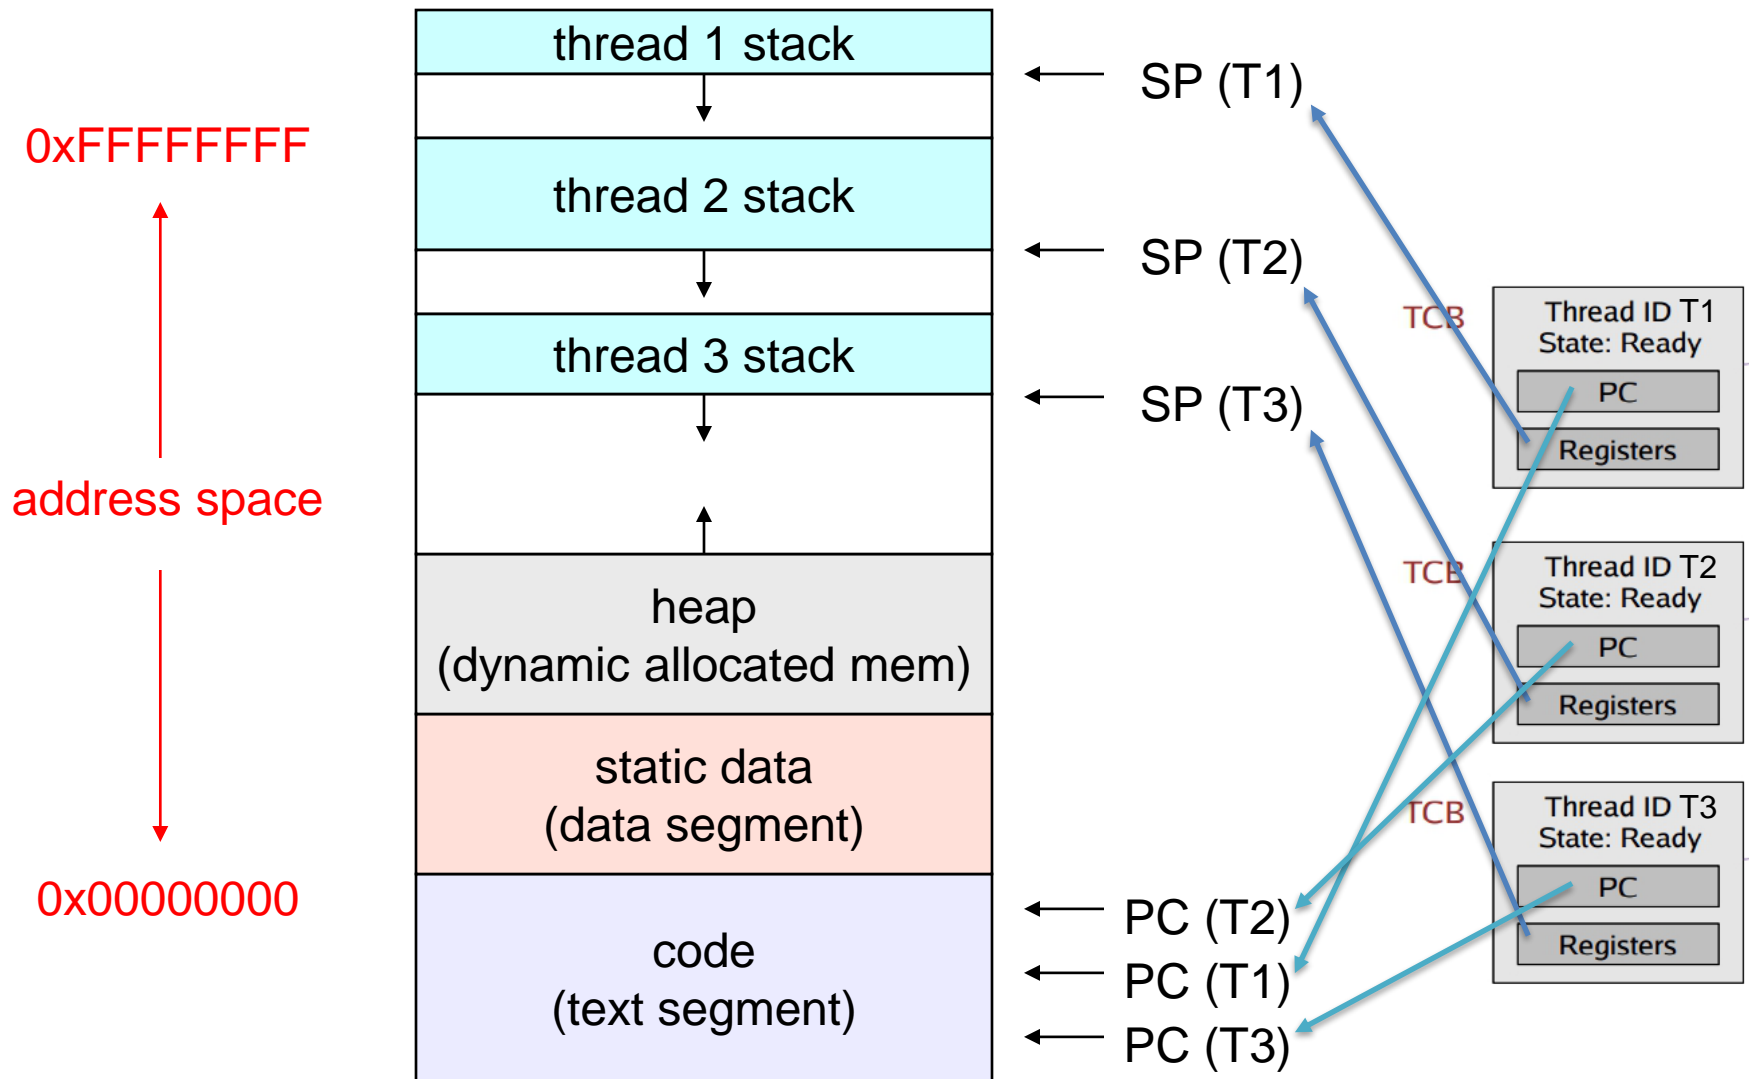
# Thread Control Block (TCB) #1

- A PCB for each process
- Break the PCB into two pieces
- Info on program execution stored in Thread Control Block (**TCB**)
  - Program counter
  - CPU registers
  - Scheduling information
  - Pending I/O information
- Other infos stored in Process Control Block (**PCB**)
  - Memory management information
  - Accounting information





# Thread Control Block (TCB) #2



# Example Applications

- Multithreading is useful for
  - Handling concurrent events (e.g., web servers and clients)
  - Building parallel programs (e.g., matrix multiply, ray tracing)
  - Improving program structure (divide and conqueror)
- Multithreading is useful **on a uniprocessor**
  - Even though only one thread can run at a time

# Terminology Note

- There is the potential for some confusion
  - “process” == “address space + OS resources + **single** execution flow”
  - **OR**
  - “process” == “address space + OS resources + **multiple** execution flows”
- Single-threaded process: 1 thread
- Multi-threaded process:  $N > 1$  threads