



THE UNIVERSITY *of* EDINBURGH
informatics

**Operating Systems
(INFR10079)
2023/2024 Semester 2**

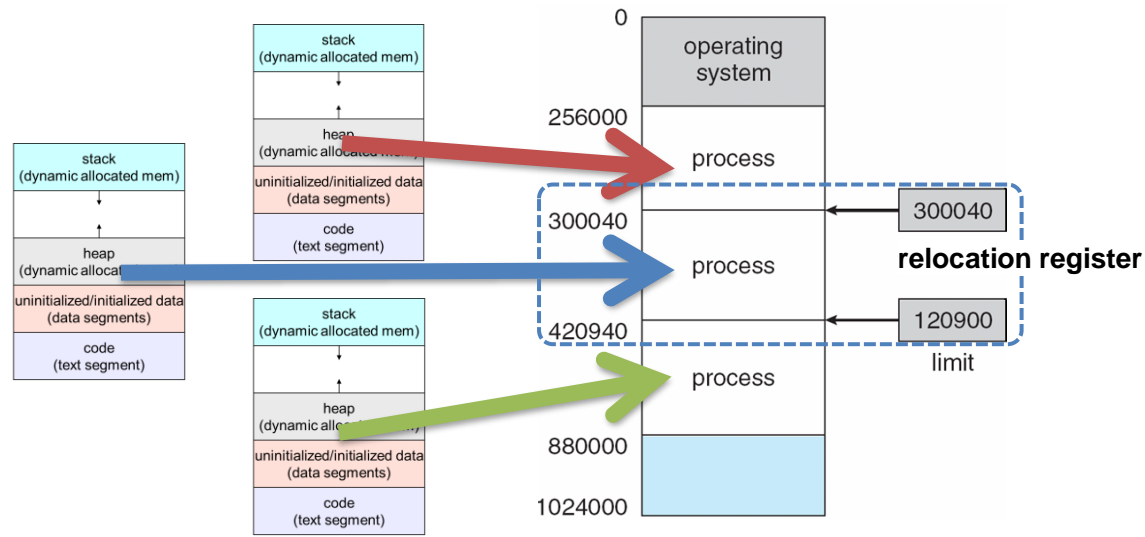
**Memory Management
(Allocation)**

abarbala@inf.ed.ac.uk

Chapter 9.1, 9.2, 9.5

Contiguous Allocation #1

- Main memory **must** support both OS and user processes
- Memory limited resource **must** allocate efficiently
- Contiguous allocation is an *early method*
- Main memory usually into two **partitions**
 - OS (usually) held in low memory with interrupt vector
 - Processes held in high memory
 - Each contained in single contiguous section of memory

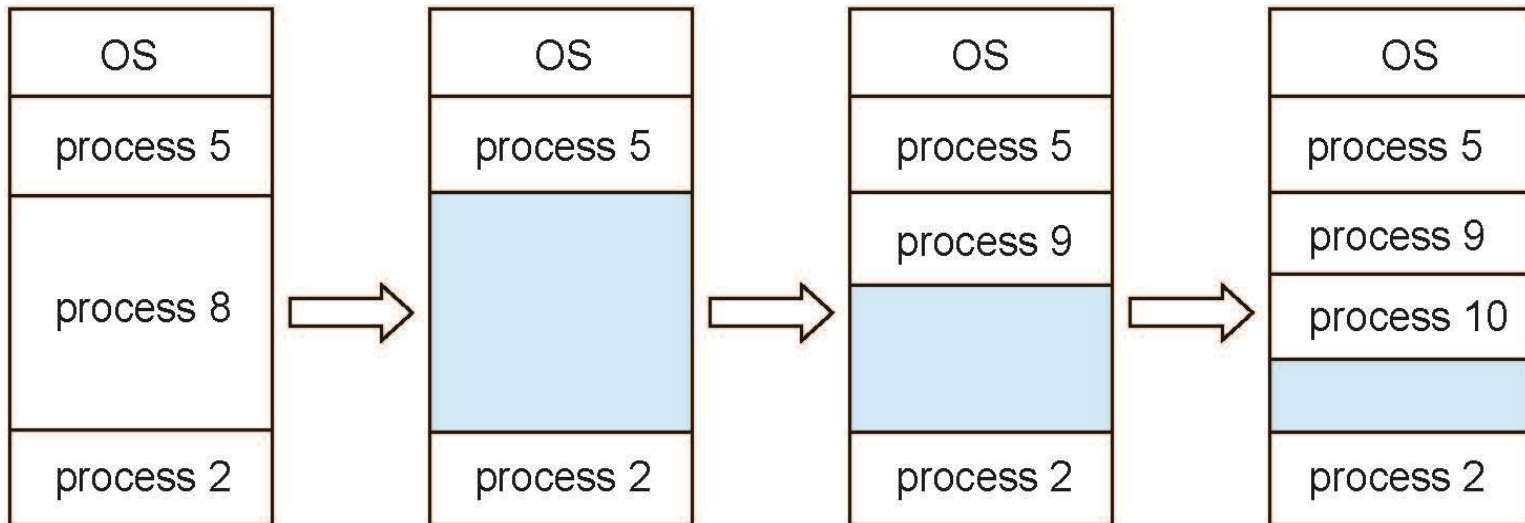


Contiguous Allocation #2

- Relocation and limit registers
 - To **protect user processes**
 - from each other
 - from changing OS code and data
 - *Relocation register* contains value of smallest physical address
 - *Limit register* contains range of logical addresses
 - Each logical address must be less than the limit register
- MMU translates logical address
 - *transparently, during execution*


Multiple-partition Allocation #1

- Multiple-partition allocation
 - Variable partition size (size of a program) **for efficiency**
 - Degree of multiprogramming **limited** by number of partitions



Multiple-partition Allocation #2

- Multiple-partition allocation
 - **Hole:** block of available memory; holes of various sizes
 - When a process arrives, **OS allocates** memory from a hole large enough to accommodate it
 - Process exiting, **returns partition to OS**, adjacent free partitions combined
 - Operating system maintains information about
 - allocated partitions
 - free partitions (hole)



Dynamic Storage-Allocation Problem

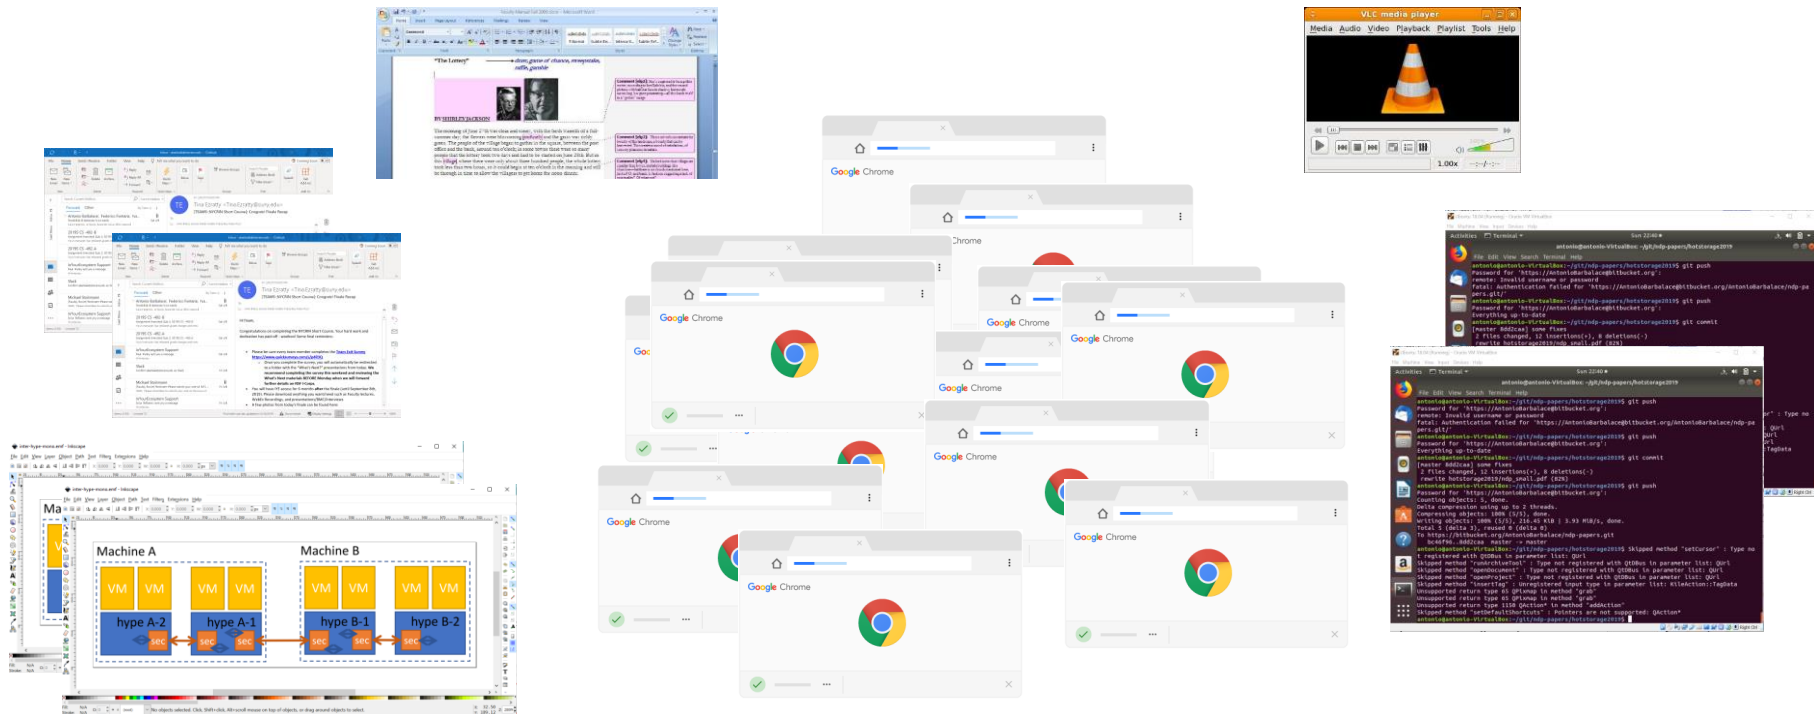
How to satisfy a request of size n from a list of free holes?

- **First-fit**: Allocate the ***first*** hole that is big enough
- **Best-fit**: Allocate the ***smallest*** hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit**: Allocate the ***largest*** hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Multiple Programs: Swapping (1)

- How many programs are currently executed?
- Do they all fit in memory?




Multiple Programs: Swapping (2)




- How many programs are currently executed?
- Do they all fit in memory?

- Swapping
 - a) **SWAP IN:** Bringing in memory a process in its entirety (from disk)
 - b) **RUN:** Running it for a while
 - c) **SWAP OUT:** Putting it back from memory (to disk)
 - About **running** processes, no programs!
 - Idle processes are stored on disk,
 - Do not take up any memory when they are not running

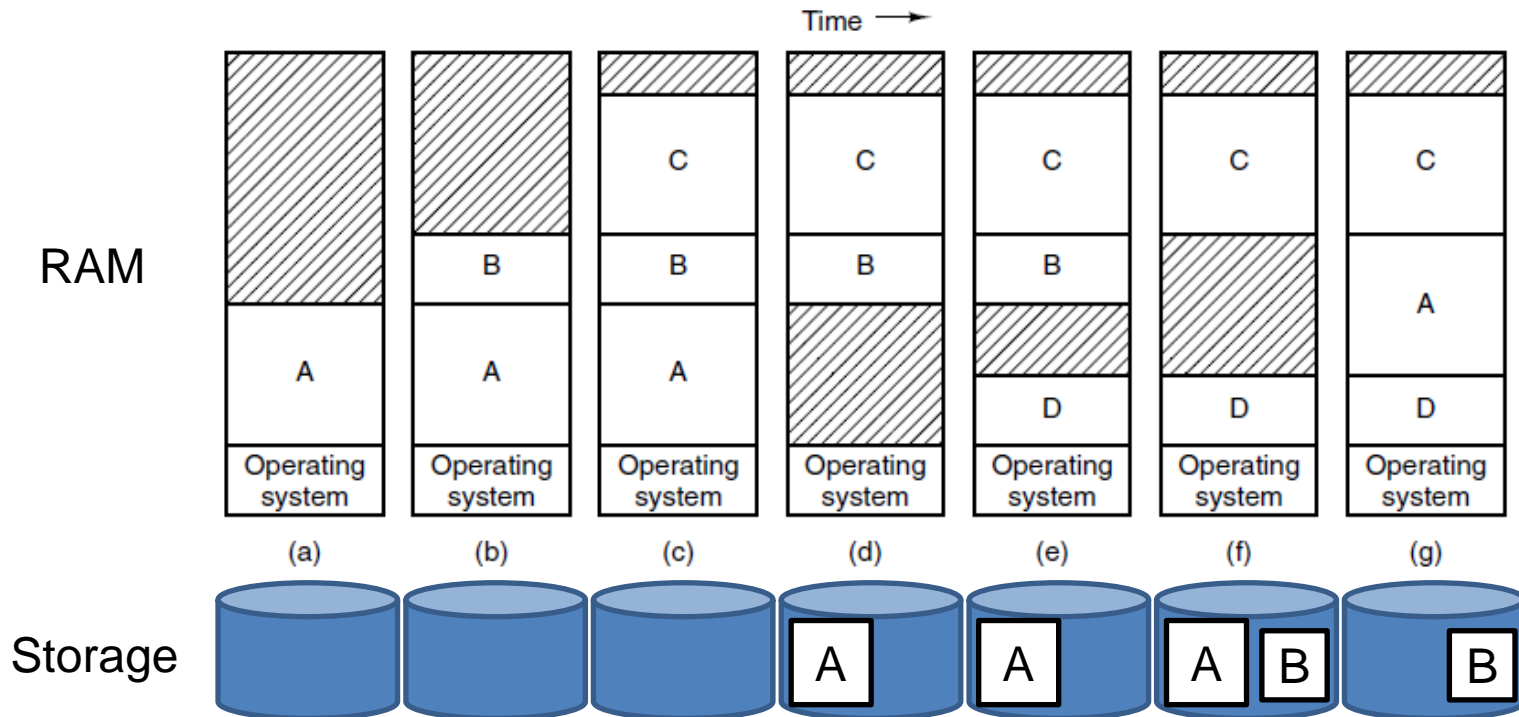


Multiple Programs: Swapping (3)




Memory allocation changes as processes come into memory and leave it. The shaded regions are unused memory (MOS Figure 3-4)

Multiple Programs: Swapping (3)

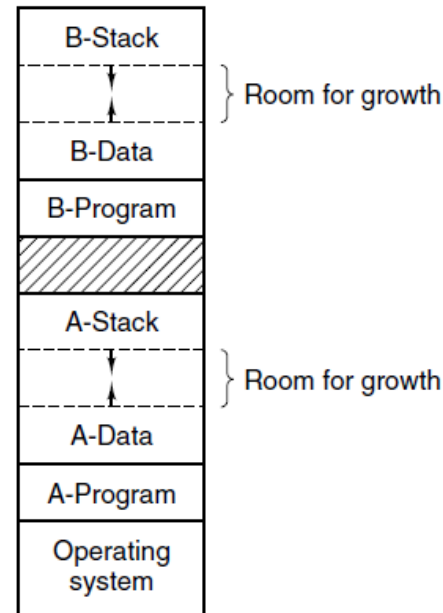


Memory allocation changes as processes come into memory and leave it. The shaded regions are unused memory (MOS Figure 3-4)

Multiple Programs: Growing Programs



(a)




(b)

- (a) Allocating space for a growing program at the end – if more space is needed, relocate
- (b) Allocating space for a growing program in its address space – max amount of growth
- (c) Both solutions are not ideal**

Multiple Programs: Memory Fragmentation #1



- Process creation and exit, swapping, process growing and shrinking
- Create **memory holes**
 - New processes, memory is available, cannot be placed
- **Compaction**
 - Move all processes tightly together
 - By memory copy or
 - By swapping out and in
 - Computationally expensive
 - What if a process needs to grow its heap or stack?



Multiple Programs: Memory Fragmentation #2



- **External Fragmentation**

- You allocate the exact amount of memory requested
- Total memory space exists to satisfy a request, but it is not contiguous

- **Internal Fragmentation**

- You allocate more than what required
- Allocated memory may be slightly larger than requested memory

- First fit analysis reveals given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable

Concepts

- Relocatable binaries
- MMU
- Logical Address Space vs Physical Address Space
- Contiguous Memory Allocation
 - First fit, best fit, worst fit
- Swapping
- Growing programs
- Fragmentation
 - Compaction