



THE UNIVERSITY of EDINBURGH
informatics

**Operating Systems
(INFR10079)
2022/2023 Semester 2**

**Memory Management
(Paging)**

abarbala@inf.ed.ac.uk

Chapter 9.3, 9.4, 9.6

Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)**
 - Points to the page table, **part of CPU**
- **Page-table length register (PTLR)**
 - Indicates size of the page table, **part of CPU**
- Every data/code access requires **two memory accesses**
 - To *fetch the **page table entry*** (translation)
 - To *fetch the actual **memory content*** (data/code)

Implementation of Page Table: TLB

- How to avoid **two** memory accesses?
- **Translation look-aside buffers (TLBs)**

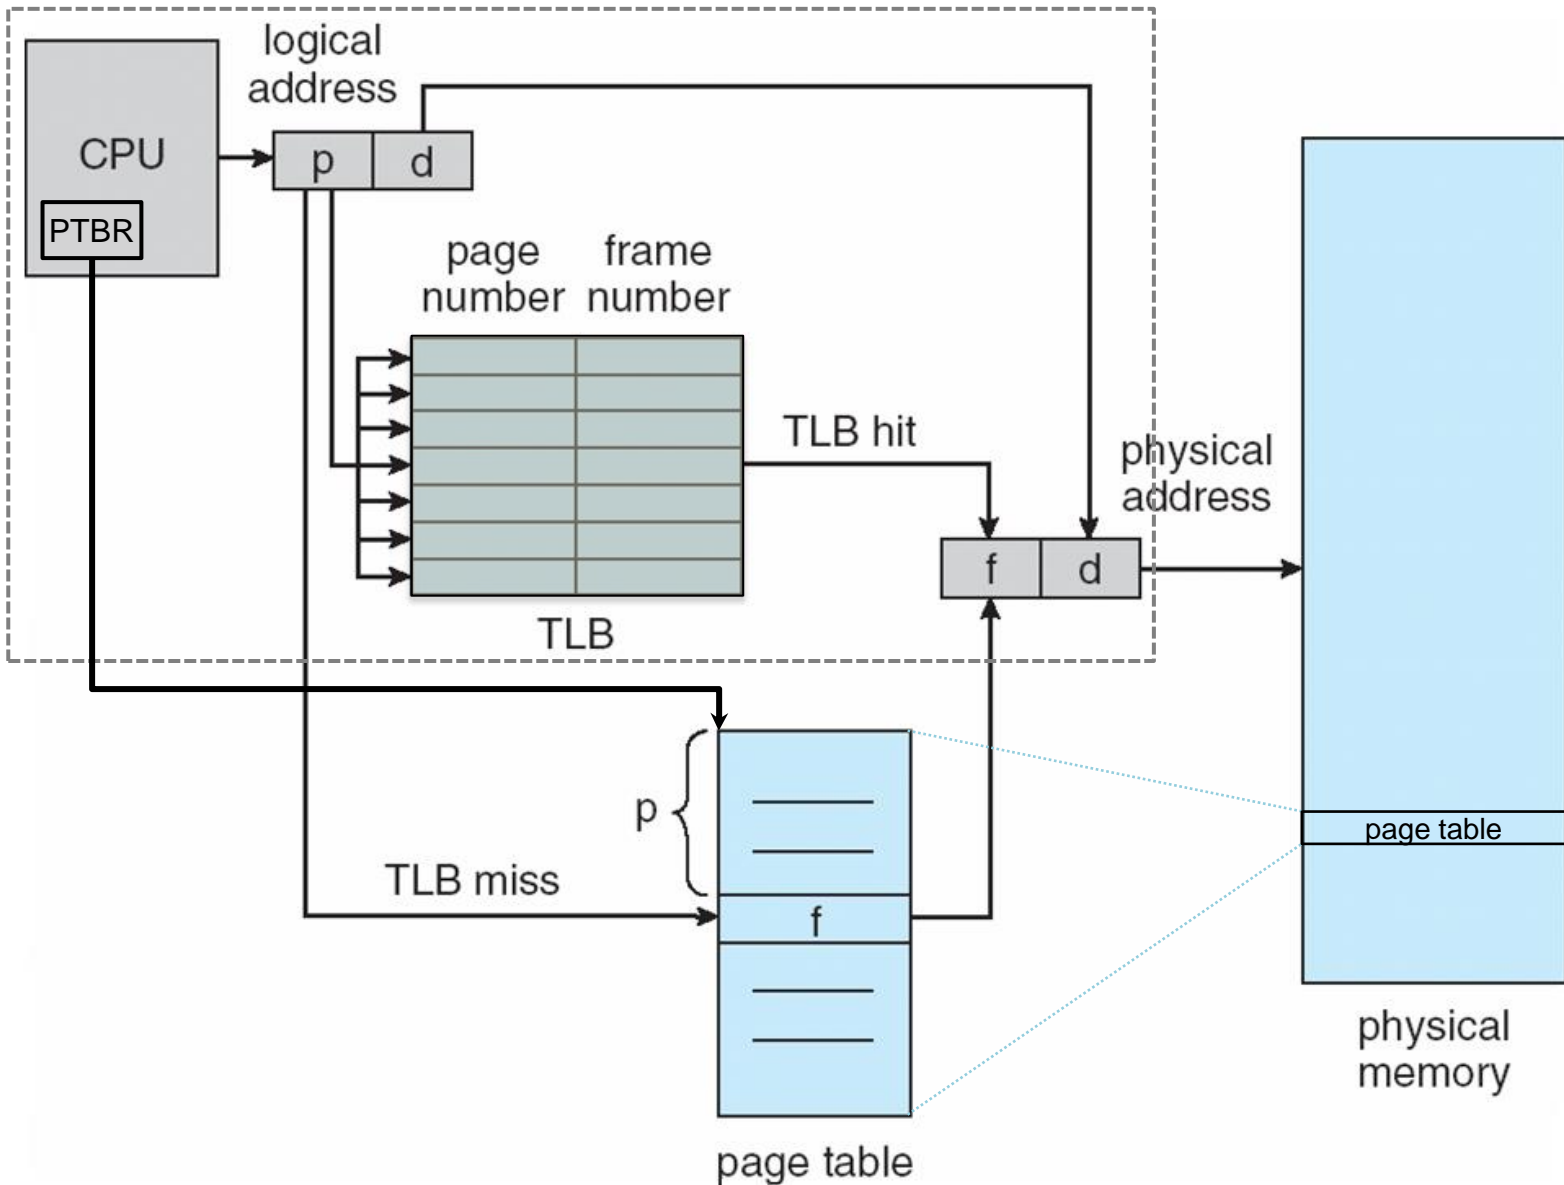
- Fast-lookup hardware cache
 - Associative memory for parallel search
 - Entry of the form **<page #, frame #>**
- Typically, small (e.g., 64 or 1,024 entries)

Page #	Frame #
0	14
2	18

- If translation
 - **Exists** in TLB (hit), use it (at no additional cost)
 - **Doesn't exist** in TLB (miss), fetch translation from memory
- Maintain translations for subsequent memory accesses
 - **Replacement policies** must be considered
 - Some entries can be **wired down** for permanent fast access

Paging Hardware With TLB

Inside CPU



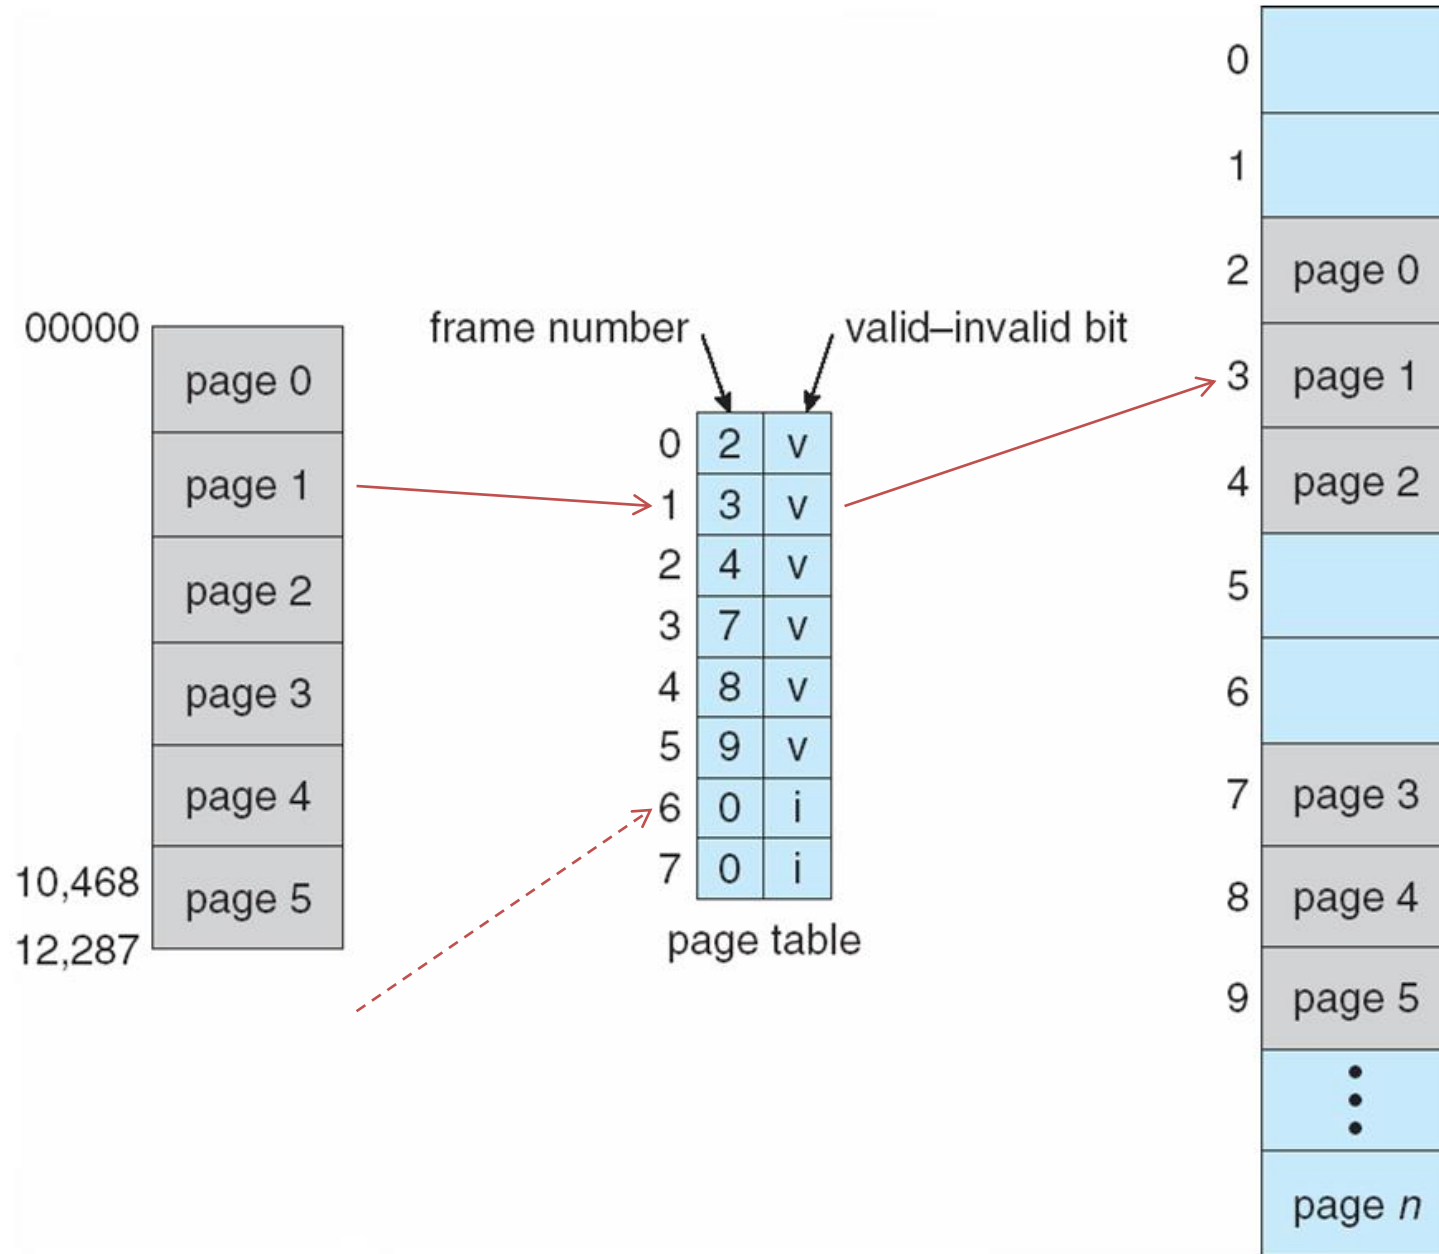
Effective Access Time (EAT) with TLB

- Memory access time
 - Time the CPU waits to access main memory directly
- Hit ratio = α
 - Percentage of times that a page number is found in TLB
- Miss ratio = $1 - \alpha$
- Consider $\alpha = 80\%$, 100ns for memory access
 - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider $\alpha = 99\%$, 100ns for memory access
 - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

Memory Protection

- **Page Table Entries (PTEs)** contain more information
- Memory protection by **protection bits** in each PTE
 - What **type of access** is allowed?
 - Read-only, read-write, execute-only
 - Does a **translation exist**?
 - **Valid** indicates that the associated page
 - **is in the process' logical address space**, thus a legal page
 - **Invalid** indicates that the page
 - **is not in the process' logical address space**
 - ...
- Any violations result in a trap to the OS kernel

Memory Protection: **Valid (v)** or **Invalid (i)** Bits



Shared Pages

- **Shared code or data**

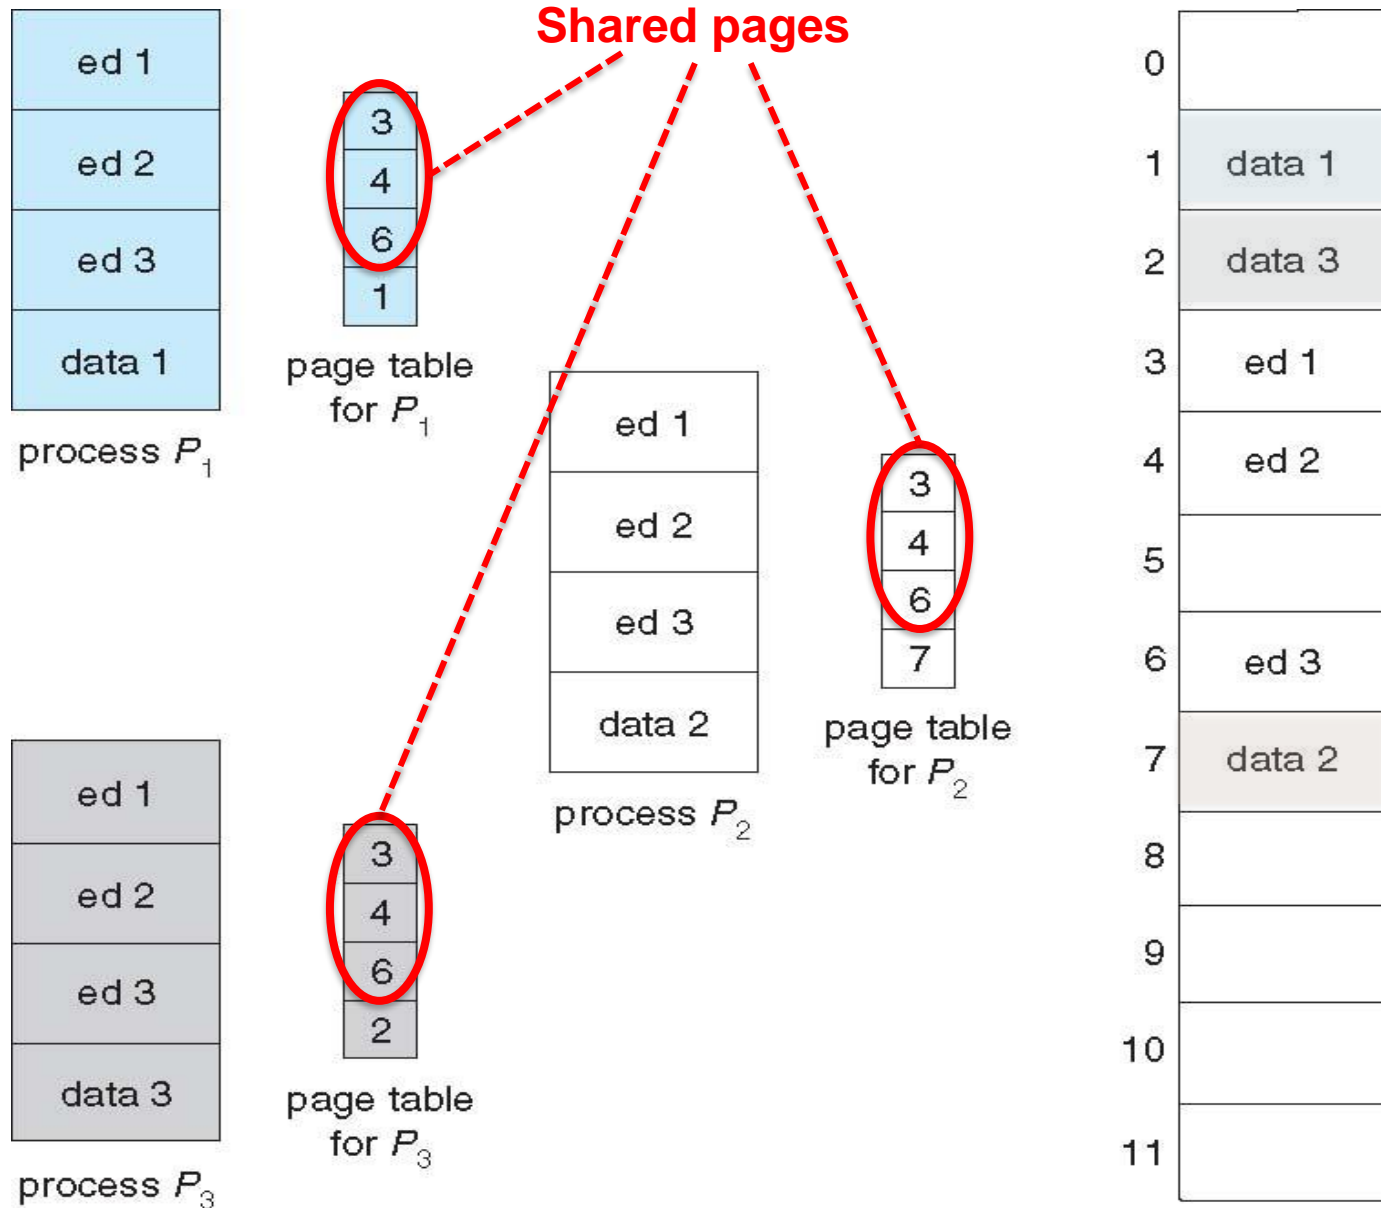
- **One copy** of read-only (reentrant) **code** shared among processes
 - Different processes may share the same library code
 - But they will have per-process data
- Read-write **data pages** shared among processes for **communication**
 - Inter-process communication
- **Somewhat different from multithreaded processes**
 - Different threads share the entire address space, and OS resources

- **Private code and data**

- Each process keeps a separate copy of the code and/or data

- Pages for the private/shared code and data can appear **anywhere** in the logical address space

Shared Pages Example




Structure of the Page Table

- Page table can **get huge**
 - Consider a 32-bit logical address space
 - Page size of 4 kB (2^{12})
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - Each entry is 4 bytes, 4 MB of memory space required
 - Only for the page table, **independently of translations**
 - Need to allocate **contiguously** in physical memory
 - Costs a lot for small memories
 - 0.1% of 4GB per process
- Alternative **constructions**
 - **Hierarchical** Page Tables
 - **Hashed** Page Tables
 - **Inverted** Page Tables


Hierarchical Page Tables

- Break up the entire logical address space into **multiple** page tables
- Then construct **another table** that refers to each of such page tables
- **Two-level page table**



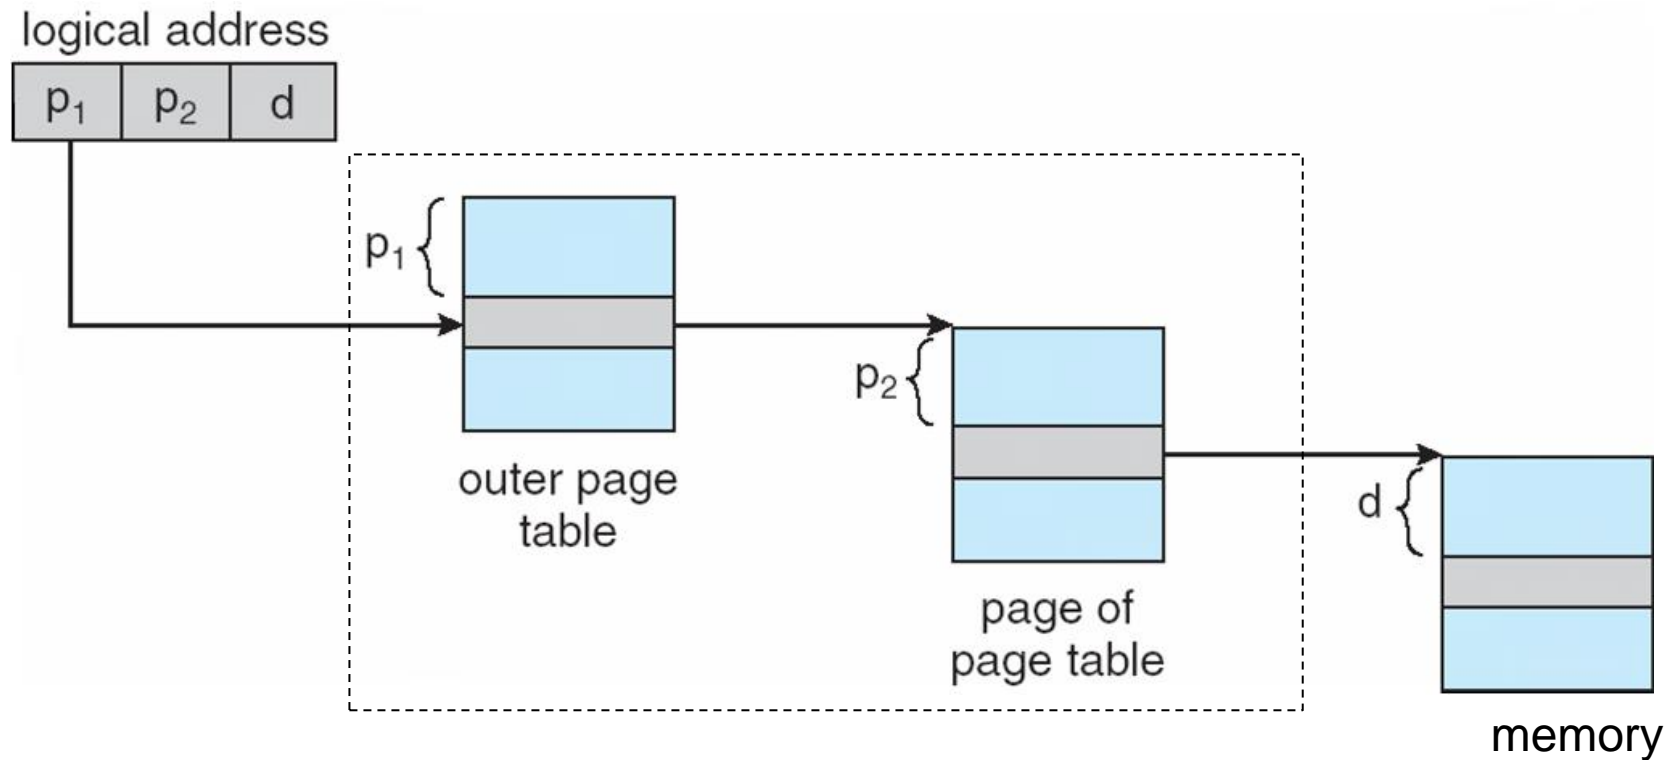
Two-Level Paging Example

- Logical address (32-bit machine with 1K page size) divided into
 - a **page number** consisting of 22 bits
 - a **page offset** consisting of 10 bits
- Page table is paged, the page number is further divided into
 - 12-bit **outer** page number
 - 10-bit **inner** page number
- A logical address looks like




- p_1 is an index into the outer page table
- p_2 is an index within the page of the inner page table

Two-Level Paging Address-Translation Scheme




64-bit Logical Address Space

- How many levels?
- If page size is 4 kB (2^{12}) and **single-level** page table
 - Page table has 2^{52} entries
- If **two-level** scheme
 - Inner page tables could be 2^{10} entries
 - Outer page table has 2^{42} entries (4,398,046,511,104)



- If **three-level** scheme – add a 2nd outer page table
 - The 2nd outer page table is still 2^{32} entries (4,294,967,296)



Multi-level Paging Scheme

- On 64bit machines logical address space $< 2^{64}$
 - For practical reasons (no machine with 18 zetta bytes)
- But 4-level and 5-level page tables exist

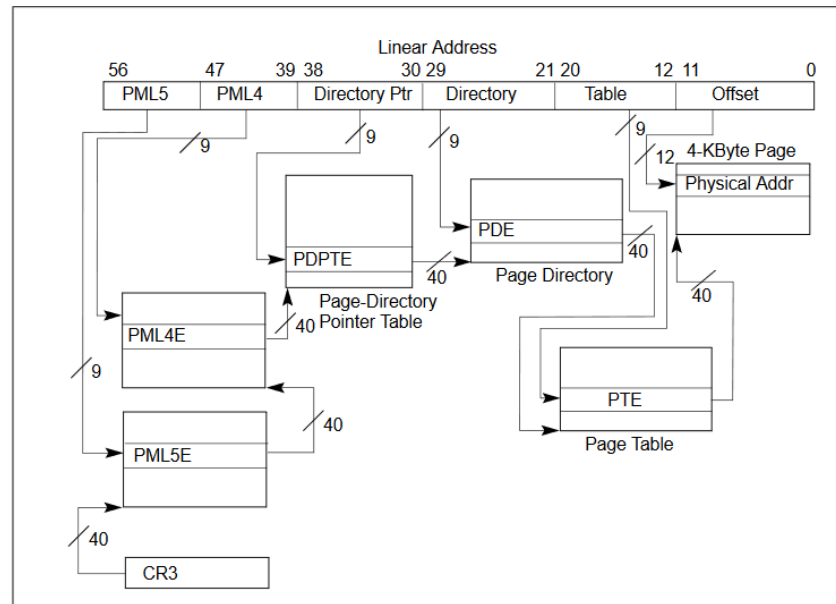



Figure 2-1. Linear-Address Translation Using 5-Level Paging

Hashed Page Tables

- The **logical page number** is hashed into a page table index
- Each entry **chains** elements hashing to the same location
- Each element **contains**
 - Logical page number
 - Value of the mapped page frame
 - Pointer to the next element
- Logical page numbers are compared searching for a **match**
 - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
 - Similar to hashed but each entry refers to several pages
 - 16 rather than 1


Hashed Page Table Example



Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages
 - **track all physical pages**
- One entry for each physical page of memory
- Entry consists of
 - **Virtual address** of the page stored in physical memory location
 - Information about the **process that owns** the page (protection)
- **Decreases memory** needed to store translations
 - Increases time needed to search the table
 - Use hash table to limit the search to one/few page-table entries
- How to implement shared memory?
 - Need OS intervention

Inverted Page Table Architecture



Why Paging? #1

- Inter-process memory protection
 - Process A address XYZ is different from Process B XYZ
- Protect code from being rewritten
 - Code pages read-only
 - A bad pointer can't change the program code
- Detect null pointer dereferencing at runtime
 - First page of the logical address space invalid
 - References to address 0 cause trap to OS
- Reduce memory usage (shared libraries)
 - Have one copy of a library in physical memory, not one per process
 - All page table entries to library point to the same set of physical frames
- Generalize use of “shared memory”
 - Regions of two processes' address spaces map to the same frames

Why Paging? #2

- **Copy-on-write (CoW)**, e.g., on fork()
 - Instead of copying all pages, create shared mappings of parent pages in child address space
 - Make shared mappings read-only for both processes
 - When either process writes, fault occurs, OS “splits” the page
- **Memory-mapped files**
 - Instead of using open, read, write, close
 - “map” a file into a region of the logical address space
 - e.g., into region with base ‘X’
 - Accessing logical address ‘X+N’ refers to offset ‘N’ in file
 - Initially, all pages in mapped region marked as invalid
 - OS reads a page from file whenever invalid page accessed
 - OS writes a page to file when evicted from physical memory
 - Only necessary if page is dirty

Summary

- Non-contiguous (physical) memory allocation
 - Reduced or no fragmentation
 - Enable sharing
 - Require hardware support
- Segmentation
 - Variable size chunks
- Paging
 - Fixed size chunks
 - Multiple optimizations