



THE UNIVERSITY *of* EDINBURGH  
**informatics**

# Operating Systems (INFR10079) 2023/2024 Semester 2

## Deadlocks


[abarbala@inf.ed.ac.uk](mailto:abarbala@inf.ed.ac.uk)



**No one can proceed** (by following its own *official* direction)


# Dining Philosophers Problem

- Dijkstra, 1965
- 5 philosophers, each with
  - A spaghetti dish
  - A fork
- To eat, a philosopher needs **two** forks
  - Spaghetti are too slippery
  - (Yes, we do that in Italy too)
- A philosopher alternates periods of
  - Eating
  - Thinking



# Deadlock Problem

- When
- all **grab** the fork at their **left** at the same time
- and **wait** for the one at their **right**
  - They will all **indefinitely wait** and no one will eat
- **Cause**
  - Each philosopher needing what another philosopher has
    - Result of sharing resources



# Deadlock Example #1

From real-life example to code example

Semaphore `mutexA = 1` /\* protects resource A \*/


Semaphore `mutexB = 1` /\* protects resource B \*/

**Process P:**

```
{  
  /* initial compute */  
  down(mutexA)  
  down(mutexB)  
  /* use both resources */  
  up(mutexA)  
  up(mutexB)  
}
```

**Process Q:**

```
{  
  /* initial compute */  
  down(mutexB)  
  down(mutexA)  
  /* use both resources */  
  up(mutexB)  
  up(mutexA)  
}
```



## Possible Deadlock



# Deadlock Example #2

From real-life example to code example

Semaphore `mutexA = 1` /\* protects resource A \*/

Semaphore `mutexB = 1` /\* protects resource B \*/




Process P:

```
{  
  /* initial compute */  
  down(mutexA)  
  4 down(mutexB) 2  
  /* use both resources */  
  up(mutexA)  
  up(mutexB)  
}
```

Process Q:

```
{  
  /* initial compute */  
  down(mutexB)  
  3 down(mutexA)  
  /* use both resources */  
  up(mutexB)  
  up(mutexA)  
}
```

## Actual Deadlock



# Deadlock Example – Try it out!

```
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;
int main() {
    pthread_mutex_init(&first_mutex, NULL);
    pthread_mutex_init(&second_mutex, NULL);
    /** * create thread one and thread two */
    /** * wait threads to finish */
    return 0;
}
/* thread one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}
/* thread two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);
    pthread_exit(0);
}
```


- Complete the program
  - Add `pthread_create`, and `pthread_join`
  - Add `printf` to track program progresses
- Compile the program
  - Don't forget `-pthread`
- Run the program
  - How many times should run it before a deadlock?

# Deadlock

- In a **multiprogramming** environment
  - Threads or processes compete for a finite number of resources
  - Each thread or process requests resources
  - If resources are not available, the thread or process waits
  - Waiting thread or process **never runs again** because the requested resource(s) is(are) held by other waiting threads or processes
- **Liveness failure**
- **Deadlock Definition**
  - A situation in which every process/thread in a set of processes/threads
  - **is waiting for an event** that can be **caused only by another process/thread in the set**



# Livelock



- Example
  - **Two people** (red/green) attempt to pass in a hallway
  - One moves to his right, the other to her left
    - still obstructing each other's progress
  - Then he moves to his left, and she moves to her right, and so forth
- They aren't blocked, but they aren't making any progress
- **Liveness failure**, similar to deadlock
  - Both prevent two or more threads/processes from proceeding
  - Threads/processes are unable to proceed **for different reasons**
- **Deadlock**
  - Occurs when every thread/process in a set **is blocked waiting** for an event that can be caused only by another thread/process in the set
- **Livelock**
  - Occurs when a thread/process **continuously attempts** an action that fails

**The next slides focus on deadlock**

# Necessary Conditions for **Deadlock**

## **A. Mutual exclusion**

- Each resource is either currently assigned to exactly **one** process/thread or is available

## **B. Hold and wait**

- Processes/threads currently holding resources that were granted earlier can request new resources

## **C. No preemption**

- Resources previously granted cannot be forcibly taken away from a process/thread
  - They must be explicitly released by the process/thread holding them

## **D. Circular wait**

- There must be a circular list of two or more processes/thread
  - Each waiting for a resource held by the next member of the chain

**All conditions must hold for a deadlock to occur**

# System Model



## (Resources and Processes)

*A system consists of a **finite number of resources**, partitioned into several types/classes, to be **distributed among a number of competing threads/processes***


- **Resource** types/classes ( $m$ )
  - $R_1, R_2, \dots, R_m$
  - E.g., printers, disks, etc.
- Each **resource** type  $R_i$  has  $E_i$  instances
  - E.g., 3 printers, 5 disks, etc.
- Assume serially reusable resources
  - request -> use -> release
- **Processes** ( $n$ )
  - $P_1, P_2, \dots, P_n$
  - Instead of processes a model with **threads** can be developed

# Resource-allocation Graph

- A way to visualize **CURRENT STATE** of threads/processes
  - Like a snapshot
- A set of vertices  $V$  and a set of edges  $E$
- $V$  is partitioned into two types
  - **Process** vertices  $P = \{P_1, P_2, \dots, P_n\}$ , the set of processes
  - **Resource** vertices  $R = \{R_1, R_2, \dots, R_m\}$ , the set of resource types
    - **Not resources!**
- $E$  is partitioned into two types
  - **Request** edge – directed edge  $P_i \rightarrow R_j$ 
    - **The process is blocked waiting for that resource**
  - **Assignment** edge – directed edge  $R_j \rightarrow P_i$ 
    - **The process owns the resource**




# Resource Allocation Graph with No Cycle and No Deadlock



Graph contains **no cycles** → No thread/process is **deadlocked**


# Resource Allocation Graph with Cycles and Deadlock



If the graph **contains a cycle**, then a deadlock may exist



# Resource Allocation Graph with Cycle and No Deadlock



A cycle **is not sufficient** to imply a deadlock

# Notes

- If each resource type has **exactly one instance**
  - Then a cycle implies that a **deadlock has occurred**
- If the cycle involves only **a set of resource types**, each of which has **only a single instance**
  - Then a **deadlock has occurred**
- Each thread involved in the cycle is deadlocked
  - A cycle in the graph is both **a necessary and a sufficient condition**
- If each resource type has **several instances**
  - then a cycle doesn't necessarily imply that a deadlock has occurred
  - a cycle in the graph **is a necessary but not a sufficient condition** for the existence of deadlock

# **HOW TO DEAL WITH DEADLOCKS?**

# Handling Deadlocks

- **Ensure deadlock never occurs**
  - **Prevention** (development time)
    - Negate one of the four necessary conditions
  - **Avoidance** (runtime)
    - Each resource request is analyzed and denied if may deadlock
- **Allow deadlock to happen**
  - **Detection and recovery**
    - System enters a deadlocked state, detect it, and recover
    - **Example:** databases
  - **Do nothing** (Ostrich algorithm)<sup>\*\*\*</sup>
    - Ignore the problem and pretend deadlocks never occur
    - **Example:** most operating systems, including Linux and Windows

# Deadlock Prevention

- Principle (Havender, 1968)
  - Ensure that at least one of the **necessary conditions** cannot hold
    - Mutual exclusion
    - Hold and wait
    - No preemption
    - Circular wait
  - Hence, deadlocks will be structurally impossible
- **Requires to change the way software is written**
  - Thus, at development time

# Attacking the **Mutual Exclusion** Condition

- **Idea**
  - Avoid assigning a resource unless absolutely necessary
  - Try to make sure that as few processes as possible may actually claim the resource
- **Practically**
  - For read, don't assign resources to a single process
    - Make data read only
  - For write, use a **mediator** which serializes the write
    - Spooler example, the mediator is the spooler daemon
- **Problem**
  - Processes may deadlock filling up the mediator



# Attacking the **Hold and Wait** Condition

- **Idea**

- Prevent processes that hold resources from waiting for more resources

- **Practically**

- Require all processes to request all their resources **before** “starting execution”
- If everything is available, the process will be allocated whatever it needs and can run to completion
- If one or more resources are busy, nothing will be allocated, and the process will just wait
  - And then reallocate everything again

- **Problem**

- Processes don't know about all needed resources


# Attacking the No Preemption Condition

- **Idea**
  - If a process is holding some resources
  - and request another resource that is not available,
  - then all resources the process is currently holding are preempted (released)
- **Practically**
  - Applied to resources whose state can be **saved and restored** later
    - CPU registers, memory space, etc.
- **Problem**
  - Not all resources can be virtualized

# Attacking the **Circular Wait** Condition

- **Idea**
  - All requests of resources must be made in numerical order by processes
- Practically
  - See Figure
- **Problem**
  - May be impossible to find an ordering to satisfy everyone

1. Imagesetter
2. Scanner
3. Plotter
4. Tape drive
5. CD Rom drive



# Deadlock Example

How to solve this with **Deadlock Prevention**

Semaphore muxA = 1 /\* protects resource A \*/

Semaphore muxB = 1 /\* protects resource B \*/

**Process P:**


```
{  
  /* initial compute */  
  down(muxA)  
  down(muxB)  
  /* use both resources */  
  up(muxB)  
  up(muxA)  
}
```

**Process Q:**

```
{  
  /* initial compute */  
  down(muxB)  
  down(muxA)  
  /* use both resources */  
  up(muxB)  
  up(muxA)  
}
```

**Process Q:**

```
{  
  /* initial compute */  
  down(muxA)  
  down(muxB)  
  /* use both resources */  
  up(muxB)  
  up(muxA)  
}
```



**Deadlock-free code!**

# Deadlock Avoidance

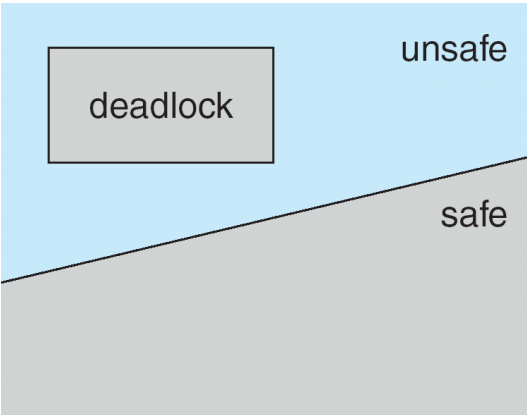
- **The OS/runtime**
- Be given **additional information** in advance
  - What resources a thread/process will **request and use** during lifetime
  - Complete **sequence of requests and releases** of each thread/process
- **Decides for each request** if a thread/process waits, considering the **resource-allocation state**
  - **Currently available** resources
  - **Currently allocated** resources to each thread/process
  - **Future requests** and releases of each thread/process
- Algorithm dynamically examines resource-allocation state
  - For **circular-wait condition** existence

# Safe State

- If the system **can allocate resources** to each thread/process (up to its maximum) **in some order** and still **avoid deadlocks**
- A system is in a safe state only if there exists a **safe sequence**
  - A sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of **ALL the processes in the systems**
  - For each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by
    - Currently **available resources** plus
    - **Resources held** by all the  $P_j$ , with  $j < i$
  - In this situation, if  $P_i$  resource needs **are not immediately available**
    - $P_i$  can wait until all  $P_j$  have finished
    - $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
    - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on



# Safe, Unsafe, Deadlock States

- A **safe state** is not a **deadlocked state**
  - A **deadlocked state** is an **unsafe state**
    - Not all **unsafe states** are **deadlocks**
  - An **unsafe state** may lead to a **deadlock**
- 
- In a **safe state**
    - OS/runtime **can avoid unsafe (and deadlocked) states**
  - In an **unsafe state**
    - OS/runtime **cannot prevent** threads from requesting resources in such a way that a **deadlock occurs**
  - Deadlock avoidance **algorithms**
    - **One instance** of each resource type: graph algorithm + claim edge
    - **Multiple instances** of each resource type: Banker algorithm


# Banker's Algorithm

- Algorithm could be used in a **banking system**
  - To ensure the bank never allocated its available cash in such a way that it could no longer satisfy the needs of all its customers
- Background
  - Set of **controlled resources** is known to the system
  - **Number of units** of each resource is known to the system
  - Each application declares **maximum possible requirement** of each resource type
- When an **application requests** a set of resource
  - System determine whether the allocation of these resources will **leave the system in a safe state**
    - YES, the resources are allocated
    - NO, **must wait until** some other releases enough resources


# Banker's Example

## Request 1

- I request a pot



- Suppose we allocate, then everyone requests its max




First, I acquire all my resources and then you do the same

**Safe state**


# Banker's Example

## Request 2

- You request a pot



- Suppose we allocate, then everyone requests its max




First, I acquire all my resources and then you do the same

**Safe state**


# Banker's Example

## Request 3a

- You request a pan



- Suppose we allocate, then everyone requests its max




First, I acquire all my resources and then you do the same (A) – not possible!  
First, you acquire all your resources and then I do the same (B) – not possible!

**UNSAFE state**


# Banker's Example

## Request 3b

- I request a pan



- Suppose we allocate, then everyone requests its max



First, I acquire all my resources and then you do the same

**Safe state**



# Banker's Algorithm Implementation #1

- Data Structures
  - $n$  = number of threads/processes
  - $m$  = number of resource types
  - **Available** – vector of length  $m$ 
    - Number of available resources of each type
  - **Max** – matrix  $n \times m$ 
    - Maximum demand of each thread/process
  - **Allocation** – matrix  $n \times m$ 
    - Number of resources of each type currently allocated to each thread or process
  - **Need** – matrix  $n \times m$ 
    - Remaining resource need of each thread/processes
  - **Need**  $[i,j] = \text{Max} [i,j] - \text{Allocation} [i,j]$ 
    - Either Need or Max are used

## Banker's Algorithm Implementation #2

- **Safety Algorithm**

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively. Initialize:  
**Work** = **Available**  
**Finish** [ $i$ ] = **false** for  $i = 0, 1, \dots, n-1$
2. Find an  $i$  such that both:  
(a) **Finish** [ $i$ ] = **false**  
(b) **Need** $_i \leq$  **Work**  
If no such  $i$  exists, go to step 4
3. **Work** = **Work** + **Allocation** $_i$   
**Finish** [ $i$ ] = **true**  
go to step 2
4. If **Finish** [ $i$ ] == **true** for all  $i$ , then the system is in a safe

# Banker's Algorithm Implementation #3

- **Resource-Request Algorithm for Process  $P_i$**

**$Request_i$**  = request vector for process  $P_i$ . If  **$Request_i[j] = k$**  then process  $P_i$  wants  **$k$**  instances of resource type  $R_j$

1. If  **$Request_i \leq Need_i$** , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  **$Request_i \leq Available$** , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

**$Available = Available - Request_i;$**

**$Allocation_i = Allocation_i + Request_i;$**

**$Need_i = Need_i - Request_i;$**

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Banker's Algorithm Example

- 5 processes  $P_0$  through  $P_4$ ;  
3 resource types:  
A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	A B C	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2	7 4 3
$P_1$	2 0 0	3 2 2		1 2 2
$P_2$	3 0 2	9 0 2		6 0 0
$P_3$	2 1 1	2 2 2		0 1 1
$P_4$	0 0 2	4 3 3		4 3 1


# Deadlock Detection and Recovery

- Allow system to **enter a deadlock state**
- **Detection**
  - Algorithm that examines the state of the system to determine whether a deadlock has occurred
    - Wait-for graph (single instance)
    - Variation of the Banker Algorithm (multiple instances)
- **Recovery scheme**
  - Algorithm to recover from the deadlock


# Detection

## Single Instance of Each Resource Type

- Maintain **wait-for** graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- **Periodically searches** for cycles in the graph
  - A cycle implies a deadlock
- An algorithm to **perform search**
  - Requires an order of  $n^2$  operations
    - $n$  is the number of vertices in the graph
- **Doesn't apply** for multiple instances of each resource type
  - Variation of the banker algorithm



**Resource allocation graph**



**Wait-for graph**

# Recovery from Deadlock

- **Process/Thread Termination**

- Abort **all** deadlocked processes
  - Abort **one** process at a time until the deadlock cycle is eliminated
    - What process to select?
- May leave resources in an **incorrect state**

- **Resource Preemption**

- Successively preempt resources from thread/processes and give them to others until deadlock cycle is broken
  1. **Selecting a victim** – resource and process to preempt
  2. **Rollback** – return to some safe state, restart process for that state
  3. **Starvation** – avoid same process/thread as victim
- Instead of recovering the **operator can be informed**