

Informatics 2D: Reasoning and Agents

Alex Lascarides

 School of
informatics



Lecture 16: Introduction to Planning

Where are we?

The first two blocks of the course dealt with ...

- Basic notions of agency
- Intelligent problem-solving
- Heuristic search, constraints
- Logic & logical reasoning
- Reasoning about actions and time

In the remainder of the course we will talk about ...

- Planning
- Uncertainty

Where are we?

The first two blocks of the course dealt with ...

- Basic notions of agency
- Intelligent problem-solving
- Heuristic search, constraints
- Logic & logical reasoning
- Reasoning about actions and time

In the remainder of the course we will talk about ...

- Planning
- Uncertainty

What is planning?

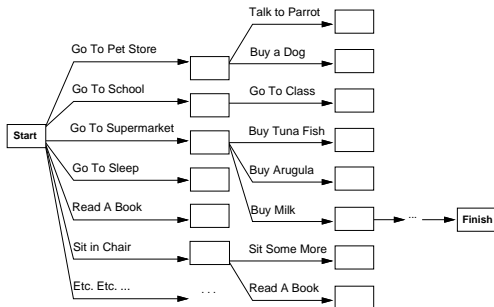
- **Planning** is the task of coming up with a sequence of actions that will achieve a goal
- We are only considering **classical planning** in which environments are
 - fully observable (accessible),
 - deterministic,
 - finite,
 - static (up to agents' actions),
 - discrete (in actions, states, objects and events).
- (Lifting some of these assumptions will be the subject of the “uncertainty” part of the course)

Why planning?

- So far we have dealt with two types of agents:
 - ① Search-based problem-solving agents
 - ② Logical planning agents
- Do these techniques work for solving planning problems?

Why planning?

- Consider a search-based problem-solving agent in a robot shopping world
- Task: Go to the supermarket and get milk, bananas and a cordless drill
- What would a search-based agent do?



Problems with search

- No goal-directedness.
- No problem decomposition into sub-goals that build on each other
 - May undo past achievements
 - May go to the store 3 times!
- Simple goal test doesn't allow for the identification of milestones
- How do we find a good heuristic function?
How do we model the way humans perceive complex goals and the quality of a plan?

How about logic & deductive inference?

- Generally a good idea, allows for “opening up” representations of states, actions, goals and plans
- If $Goal = Have(Bananas) \wedge Have(Milk)$ this allows achievement of sub-goals (if independent)
- Current state can be described by properties in a compact way (e.g. $Have(Drill)$ stands for hundreds of states)
- Allows for compact description of actions, for example

$$Object(x) \Rightarrow Can(a, Grab(x))$$


- Allows for representing a plan hierarchically, e.g. $GoTo(Supermarket) = Leave(House) \wedge ReachLocationOf(Supermarket) \wedge Enter(Supermarket)$ then decompose further into sub-plans

How about logic & deductive inference?


Problems:

- 1 In its general form either awkward (propositional logic) or tractability problems (first-order logic)
- 2 If p is a sequence that achieves the goal, then so is $[a, a^{-1}|p]!$
- 3 (Logically independent) subgoals may need to be undone to achieve other goals.

Goal: $on(A, B) \wedge on(B, C)$



Initial State



Goal

What next?

Solutions: We need

- 1 To reduce complexity to allow scaling up.
- 2 To allow reasoning to be guided by plan 'quality'/efficiency.

Do 1. next, and 2. after that.

Representing planning problems

- Need a language expressive enough to cover interesting problems, restrictive enough to allow efficient algorithms.
- **Planning Domain Definition Language** or **PDDL**
- PDDL will allow you to express:
 - 1 states
 - 2 actions: a description of transitions between states
 - 3 and goals: a (partial) description of a state.

Representing States and Goals in PDDL

- **States** represented as conjunctions of propositional or function-free first order positive literals:
 - $Happy \wedge Sunshine,$
 $At(Plane_1, Melbourne) \wedge At(Plane_2, Sydney)$
- So these **aren't states**:
 - $At(x, y)$ (no variables allowed),
 $Love(Father(Fred), Fred)$ (no function symbols allowed)
 $\neg Happy$ (no negation allowed).

Closed-world assumption!

- A **goal** is a **partial description** of a state, and you can use negation, variables etc. to express that description.
 - $\neg Happy, At(x, SFO), Love(Father(Fred), Fred) \dots$

Actions in PDDL

Action(*Fly*(*p*, *from*, *to*),
Precond: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$
Effect: $\neg At(p, from) \wedge At(p, to)$)

- Actually **action schemata**, as they may contain variables
- Action name and parameter list serves to identify the action
- **Precondition**: defines states in which action is **executable**:
 - Conjunction of positive and negative literals, where all variables must occur in action name.
- **Effect**: defines how literals in the input state get changed (anything not mentioned stays the same).
 - Conjunction of positive and negative literals, with all its variables also in the preconditions.
 - Often positive and negative effects are divided into **add list**

The semantics of PDDL: States and their Descriptions

- $s \models At(P_1, SFO)$ iff $At(P_1, SFO) \in s$
 $s \models \neg At(P_1, SFO)$ iff $At(P_1, SFO) \notin s$
 $s \models \phi(x)$ iff there is a ground term d such that $s \models \phi[x/d]$.
 $s \models \phi \wedge \psi$ iff $s \models \phi$ and $s \models \psi$

The Semantics of PDDL: Applicable Actions

- Any action is **applicable** in any state that satisfies the precondition with an appropriate substitution for parameters.
- Example: State

$$\begin{aligned} &At(P_1, Melbourne) \wedge At(P_2, Sydney) \wedge Plane(P_1) \wedge Plane(P_2) \\ &\wedge Airport(Sydney) \wedge Airport(Melbourne) \wedge Airport(Heathrow) \end{aligned}$$

satisfies

$$At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$$

with substitution (among others)

$$\{p/P_2, from/Sydney, to/Heathrow\}$$

The semantics of PDDL: The Result of an Action

- **Result** of executing action a in state s is state s' with any positive literal P in a 's Effects added to the state and every negative literal $\neg P$ removed from it (under the given substitution) .
- In our example s' would be

$$\begin{aligned} &At(P_1, Melbourne) \wedge At(P_2, Heathrow) \wedge Plane(P_1) \wedge Plane(P_2) \\ &\wedge Airport(Sydney) \wedge Airport(Melbourne) \wedge Airport(Heathrow) \end{aligned}$$

- “PDDL assumption”: every literal not mentioned in the effect remains unchanged (cf. frame problem)
- **Solution** = action sequence that leads from the initial state to a state that satisfies the goal.

Blocks world example

- Given: A set of cube-shaped blocks sitting on a table
- Can be stacked, but only one on top of the other
- Robot arm can move around blocks (one at a time)
- Goal: to stack blocks in a certain way
- Formalisation in PDDL:
 - $On(b, x)$ to denote that block b is on x (block/table)
 - $Move(b, x, y)$ to indicate action of moving b from x to y
 - Precondition for this action requires $Clear(z)$: nothing stacked on z .

Blocks world example

- Action schema:

Action(*Move*(*b*, *x*, *y*),

Precond: $On(b, x) \wedge Clear(b) \wedge Clear(y)$

Effect: $On(b, y) \wedge Clear(x) \wedge \neg On(b, x) \wedge \neg Clear(y)$)

- Problem: when $x = Table$ or $y = Table$ we infer that the table is clear when we have moved a block from it (not true) and require that table is clear to move something on it (not true)
- Solution: introduce another action

Action(*MoveToTable*(*b*, *x*),

Precond: $On(b, x) \wedge Clear(b)$

Effect: $On(b, Table) \wedge Clear(x) \wedge \neg On(b, x)$)

Does this Work?

- Interpret $Clear(b)$ as “there is space on b to hold a block” (thus $Clear(Table)$ is always true)
- But without further modification, planner can still use $Move(b, x, Table)$:
 - Needlessly increases search space (not a big problem here, but can be)
- So part of solution is to also add $Block(b) \wedge Block(y)$ to precondition of $Move$

Summary

- Defined the planning problem
 - Discussed problems with search/logic
 - Introduced PDDL: a special representation language for planning
 - Blocks world example as a famous application domain
 - Next time: Algorithms for planning!
- State-Space Search and Partial-Order Planning**